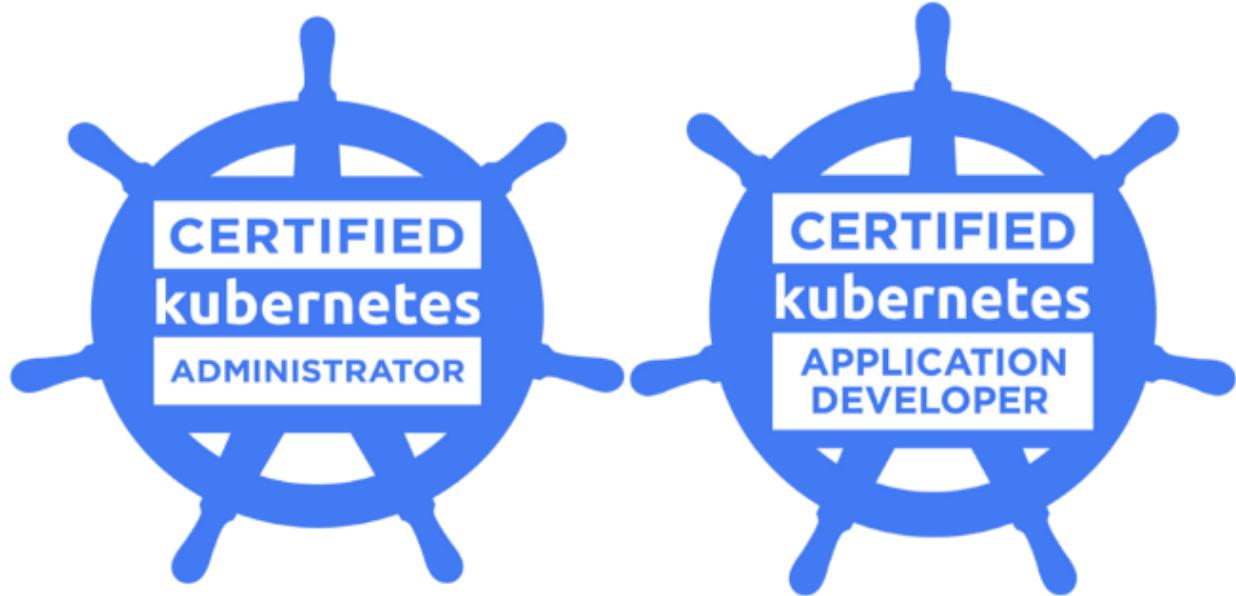


Kubernetes Master Document

(CKA & CKAD)



Content :

1. Display Commands
 - 1.1. Check the Nodes connected with Master
 - 1.2. Check Details of Nodes connected with Master
 - 1.3. Check all the Namespace in Kubernetes Node
 - 1.4. Check Pods running in specific Namespace “monitoring”
 - 1.5. Check running Pods in internal Namespace “kube-system”
 - 1.6. Check the Labels attached to Pods
 - 1.7. Describe specific Node
 - 1.8. Check cluster-info
 - 1.9. check API resources KIND name for YAML deployment

- 1.10. Check Replica Set
- 1.12. Describe Replica Set
- 1.12. List all deployment in cluster
- 1.13. Describe a specific deployment in cluster
- 1.14. Check daemonset in cluster
- 1.15. Describe daemonset in specific namespace
- 1.16. Check the events in a specific Namespace
2. Cordon and UnCordon a specific Node
3. Apply/Overwrite/Remove Label on Node
 - 3.1. Use -L to find all nodes using disk as key , value can be any
 - 3.2. Overwrite or change the value of Label on any Nodes
 - 3.3. Remove the Label disk
4. Create / delete Namespace
5. Manage cluster from Local host using kubeconfig
6. Create / delete / Rollout / Rollback Deployment
7. Delete the Stuck pods in Terminating State
8. Edit deployment to change any values like maxSurge or maxUnavailable
9. Create/Delete/Verify pods using DaemonSet apply
 - 9.1. Prepare DaemonSet yaml .
 - 9.2. Execute DaemonSet yaml using apply
 - 9.3. Check DaemonSet
 - 9.4. Describe DaemonSet
 - 9.5. Delete DaemonSet
 - 9.6. Nodeselector Scheduling
 - 9.7. nodeName Scheduling
10. Taint and Toleration

11. Affinity and AntiAffinity
12. Secrets and Config Maps
13. Resource Quota
14. Request and Limit
15. Edit any kubernetes resource with option -o yaml & Delete the pods after editing replicaset
16. Scale replicaSet
17. Creating static pod without replicaset and deployment
18. Find a Pod in all namespace
19. Creating pod using Imperative command using dr-run=client
20. Creating a ClusterIP service using kubectl expose
21. Creating deployment using Imperative command
22. Create a pod with container port 8080
23. Create ns,deployment using Imperative command
24. Create a httpd pod and expose it on port 80
25. Enter Sleep with timer in container command
26. Create ConfigMap using Command
27. Create Secret having multiple key using command
28. Find out user used to execute the sleep process within the ubuntu-sleeper pod ?
29. Change default user to run command in container using securityContext
30. save running Pod output to yaml file
31. securityContext capability
32. Multiple capabilities in Array
33. Check existing serviceaccounts
34. Create token using serviceAccount

35. Check logs of a pod
36. Liveness, Readiness, and Startup Probes
37. check memory of nodes in cluster
38. InitContainer yaml
39. Find Resources with specific Multiple labels
40. Create Job using yaml and Imperative command
41. Describe job to check success and failed logs in events
42. Job finishes in 1 sec using 3 parallel job parallelism: 3
43. Schedule cron job in yaml
44. Check the TargetPort configured in specific service svc
45. Create a NodePort svc to access the web-app from outside of cluster
46. Network Policy-networkpolicy-netpol
47. Create Policy to allow pod access from outside
48. Check Ingress Resource in all namespace
49. Different Content is shown on /wear and /watch as per Backends configured in Ingress
50. Edit the key values in Ingress using edit
51. Create a new Ingress for path /pay and backend service name : pay-service
52. Roles and RoleBindings in Service Account
53. Check Logs of Container from outside of Pod (From Host)
54. Create a Persistent Volume
55. Create a pvc Persistent Volume Claim
56. STORAGE CLASS
57. create a Pod using persistentVolumeClaim
58. switch to other context using kubectl
59. make the new my-kube-config as default config file

60. Identify the authorization modes configured on the cluster.
61. Roles
62. Which account is mapped to the Role kube-proxy
63. execute command using other user — as
64. Create a Role and roleBinding to access Pods from dev-user
65. Update the Role to create Pod using deployment .
66. Count the number of clusterroles in cluster using wc
67. Add storage access to existing role michelle
68. Add NamespaceAutoProvision in — enable-admission-plugins to allow creation of namespace while creating Pod in new namespace
69. Create a TLS secret for webhook
70. Find short names for multiple resources in api-resources
71. Find major minor in kubernetes version
72. Enable the v1alpha1 version for rbac.authorization.k8s.io API group
73. change api version using kubectl-convert plugin
74. custom Resource (customresourcedefinitions)
75. Create a new custom resource using customresourcedefinitions
76. Deployment Strategy
77. decode password from secret using base64 -d
78. upgrading from Client side to server side and vice versa

1) Display Commands

1.1) Check the Nodes connected with Master

```
kubectl get nodes
```

```
[root@master-node1 ~]# kubectl get nodes
NAME      STATUS    ROLES   AGE   VERSION
master-node1  NotReady  master  19d   v1.18.20
worker-node1  NotReady  <none>  19d   v1.18.20
worker-node2  NotReady  <none>  19d   v1.18.20
```

1.2) Check Details of Nodes connected with Master

```
kubectl get nodes -o wide
```

```
[root@master-node1 ~]# kubectl get nodes -o wide
NAME      STATUS    ROLES   AGE   VERSION   INTERNAL-IP   EXTERNAL-IP   OS-IMAGE         KERNEL-VERSION   CONTAINER-RUNTIME
master-node1  NotReady  master  19d   v1.18.20  192.168.8.136  <none>       CentOS Linux 7 (Core)  3.18.0-1168.88.1.el7.x86_64  docker://18.9.1
worker-node1  NotReady  <none>  19d   v1.18.20  192.168.8.137  <none>       CentOS Linux 7 (Core)  3.18.0-1168.88.1.el7.x86_64  docker://18.9.1
worker-node2  NotReady  <none>  19d   v1.18.20  192.168.8.126  <none>       CentOS Linux 7 (Core)  3.18.0-1168.88.1.el7.x86_64  docker://18.9.1
```

1.3) Check all the Namespace in Kubernetes Node

```
kubectl get ns
```

```
[root@master-node1 ~]# kubectl get ns
NAME      STATUS    AGE
default   Active   19d
kube-node-lease  Active   19d
kube-public  Active   19d
kube-system  Active   19d
```

1.4) Check Pods running in specific Namespace “monitoring”

```
kubectl get pods -n monitoring
```

NAME	READY	STATUS	RESTARTS	AGE
alertmanager-main-0	2/2	Running	0	92d
alertmanager-main-1	2/2	Running	0	92d
alertmanager-main-2	2/2	Running	0	92d
blackbox-exporter-64db9f4fdc-kh84d	3/3	Running	0	92d
grafana-76b545b8bd-nflrn	1/1	Running	0	92d
kube-state-metrics-cf85bcf5b-xvnxj	3/3	Running	0	92d
node-exporter-7489q	2/2	Running	0	92d
node-exporter-9fvw7	2/2	Running	0	92d
node-exporter-vdbhd	2/2	Running	0	92d
prometheus-adapter-74865dcff-dx9c2	1/1	Running	0	92d
prometheus-adapter-74865dcff-n5fz8	1/1	Running	0	92d
prometheus-k8s-0	2/2	Running	1	43d
prometheus-k8s-1	2/2	Running	1	43d
prometheus-operator-595699887b-gvjgv	2/2	Running	0	92d

1.5) Check running Pods in internal Namespace “kube-system”

```
kubectl get pods -n kube-system
```

NAME	READY	STATUS	RESTARTS	AGE
coredns-66bff467f8-9xmmh	1/1	Running	0	19d
coredns-66bff467f8-jtj5r	1/1	Running	0	19d
etcd-master-node1	1/1	Running	2	19d
kube-apiserver-master-node1	1/1	Running	2	19d
kube-controller-manager-master-node1	1/1	Running	2	19d
kube-proxy-755b4	1/1	Running	0	19d
kube-proxy-8h59m	1/1	Running	0	19d
kube-proxy-9ztlq	1/1	Running	2	19d
kube-scheduler-master-node1	1/1	Running	2	19d
weave-net-dnxk5	0/2	Init:ImageInspectError	0	19d
weave-net-l4m52	0/2	PodInitializing	0	19d
weave-net-l94nw	2/2	Running	1	19d

```
kubectl get pods -n kube-system -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
coredns-66bff467f8-9xmmh	1/1	Running	0	19d	10.32.0.3	worker-node1	<none>	<none>
coredns-66bff467f8-jtj5r	1/1	Running	0	19d	10.32.0.2	worker-node1	<none>	<none>
etcd-master-node1	1/1	Running	2	19d	192.168.0.136	master-node1	<none>	<none>
kube-apiserver-master-node1	1/1	Running	2	19d	192.168.0.136	master-node1	<none>	<none>
kube-controller-manager-master-node1	1/1	Running	2	19d	192.168.0.136	master-node1	<none>	<none>
kube-proxy-755b4	1/1	Running	0	19d	192.168.0.126	worker-node2	<none>	<none>
kube-proxy-8h59m	1/1	Running	0	19d	192.168.0.137	worker-node1	<none>	<none>
kube-proxy-9ztlq	1/1	Running	2	19d	192.168.0.136	master-node1	<none>	<none>
kube-scheduler-master-node1	1/1	Running	2	19d	192.168.0.136	master-node1	<none>	<none>
weave-net-dnxk5	0/2	Init:ImageInspectError	0	19d	192.168.0.136	master-node1	<none>	<none>
weave-net-l4m52	0/2	PodInitializing	0	19d	192.168.0.126	worker-node2	<none>	<none>
weave-net-l94nw	2/2	Running	1	19d	192.168.0.137	worker-node1	<none>	<none>

1.6) Check the Labels attached to Pods

```
kubectl get pods -n kube-system --show-labels
```

NAME	READY	STATUS	RESTARTS	AGE	LABELS
coredns-660ff467f8-9xmh	1/1	Running	0	19d	k8s-app=kube-dns,pod-template-hash=660ff467f8
coredns-660ff467f8-tj5r	1/1	Running	0	19d	k8s-app=kube-dns,pod-template-hash=660ff467f8
etcd-master-node1	1/1	Running	2	19d	component=etcd,tier=control-plane
kube-apiserver-master-node1	1/1	Running	2	19d	component=kube-apiserver,tier=control-plane
kube-controller-manager-master-node1	1/1	Running	2	19d	component=kube-controller-manager,tier=control-plane
kube-proxy-755b6	1/1	Running	0	19d	controller-revision-hash=5bd57b4bf,k8s-app=kube-proxy,pod-template-generation=1
kube-proxy-8h59m	1/1	Running	0	19d	controller-revision-hash=5bd57b4bf,k8s-app=kube-proxy,pod-template-generation=1
kube-proxy-9ztlq	1/1	Running	2	19d	controller-revision-hash=5bd57b4bf,k8s-app=kube-proxy,pod-template-generation=1
kube-scheduler-master-node1	1/1	Running	2	19d	component=kube-scheduler,tier=control-plane
weave-net-dhxk5	0/2	Init:ImageInspectError	0	19d	controller-revision-hash=7bbf1cc988, name=weave-net,pod-template-generation=1
weave-net-14652	0/2	PodInitializing	0	19d	controller-revision-hash=7bbf1cc988, name=weave-net,pod-template-generation=1
weave-net-194mw	2/2	Running	1	19d	controller-revision-hash=7bbf1cc988, name=weave-net,pod-template-generation=1

1.7) Describe specific Node

```
kubectl describe nodes worker-node1
```

[root@master-node1 ~]# kubectl describe nodes worker-node1								
Name:	worker-node1							
Roles:	<none>							
Labels:	beta.kubernetes.io/arch=amd64 beta.kubernetes.io/os=linux kubernetes.io/arch=amd64 kubernetes.io/hostname=worker-node1 kubernetes.io/os=linux							
Annotations:	kubeadm.alpha.kubernetes.io/cri-socket: /var/run/dockershim.sock node.alpha.kubernetes.io/ttl: 0 volumes.kubernetes.io/controller-managed-attach-detach: true							
CreationTimestamp:	Sun, 20 Nov 2022 14:43:35 -0500							
Taints:	node.kubernetes.io/unreachable:NoSchedule							
Unschedulable:	false							
Lease:								
HolderIdentity:	worker-node1							
AcquireTime:	<sunset>							
RenewTime:	Sun, 20 Nov 2022 15:11:41 -0500							
Conditions:								
Type	Status	LastHeartbeatTime	LastTransitionTime					
---	---	---	---					
NetworkUnavailable	False	Sun, 20 Nov 2022 14:46:02 -0500	Sun, 20 Nov 2022 14:46:02 -0500					
MemoryPressure	Unknown	Sun, 20 Nov 2022 15:08:41 -0500	Sat, 10 Dec 2022 02:33:59 -0500					
DiskPressure	Unknown	Sun, 20 Nov 2022 15:08:41 -0500	Sat, 10 Dec 2022 02:33:59 -0500					
PIDPressure	Unknown	Sun, 20 Nov 2022 15:08:41 -0500	Sat, 10 Dec 2022 02:33:59 -0500					
Ready	Unknown	Sun, 20 Nov 2022 15:08:41 -0500	Sat, 10 Dec 2022 02:33:59 -0500					
Addresses:								
InternalIP:	192.168.0.137							

1. 8) Check cluster-info

```
kubectl cluster-info
```

```
[root@master-node1 .kube]# kubectl cluster-info
Kubernetes master is running at https://192.168.0.136:6443
KubeDNS is running at https://192.168.0.136:6443/api/v1/namespaces/kube-system/services/kube-dns:dns/proxy
To further debug and diagnose cluster problems, use 'kubectl cluster-info dump'.
```

1.9) check API resources KIND name for YAML deployment

```
kubectl api-resources | less
```

```
[root@master-node1 .kube]# kubectl api-resources | less
NAME           SHORTNAMES   APIGROUP      NAMESPACED   KIND
bindings       binding      v1            true        Binding
componentstatuses   cs          v1            false       ComponentStatus
configmaps      cm          v1            true        ConfigMap
endpoints       ep          v1            true        Endpoints
events          ev          v1            true        Event
limitranges     limits       v1            true        LimitRange
namespaces      ns          v1            false       Namespace
nodes          node         v1            false       Node
persistentvolumeclaims   pvc         v1            true        PersistentVolumeClaim
persistentvolumes    pv          v1            false       PersistentVolume
pods           pod          v1            true        Pod
podtemplates    podtemplate  v1            true        PodTemplate
replicationcontrollers   rc          v1            true        ReplicationController
resourcequotas   quota        v1            true        ResourceQuota
secrets          secret       v1            true        Secret
serviceaccounts  sa          v1            true        ServiceAccount
services         svc          v1            true        Service
mutatingwebhookconfigurations   admissionregistration.k8s.io   false       MutatingWebhookConfiguration
validatingwebhookconfigurations   admissionregistration.k8s.io   false       ValidatingWebhookConfiguration
customresourcedefinitions        crd,crds    apiextensions.k8s.io   false       CustomResourceDefinition
apiservices      apiservice    apiextensions.k8s.io   false       APIService
controllerrevisions   controllerrevision   apps          true        ControllerRevision
daemonsets       daemonset    apps          true        DaemonSet
deployments      deployment   apps          true        Deployment
replicasets      replicaset   apps          true        ReplicaSet
statefulsets     statefulset  apps          true        StatefulSet
tokenreviews     tokenreview  authentication.k8s.io  false       TokenReview
localsubjectaccesreviews   localsubjectaccesreview  authorization.k8s.io  true        LocalSubjectAccessReview
selfsubjectaccesreviews   selfsubjectaccesreview  authorization.k8s.io  false       SelfSubjectAccessReview
selfsubjectrulesreviews   selfsubjectrulesreview  authorization.k8s.io  false       SelfSubjectRulesReview
subjectaccesreviews   subjectaccessreview  authorization.k8s.io  false       SubjectAccessReview
horizontalpodautoscalers   hpa         autoscaling   true        HorizontalPodAutoscaler
cronjobs         cronjob     batch        true        CronJob
jobs             job          batch        true        Job
certificatesigningrequests   csr         certificates.k8s.io  false       CertificateSigningRequest
```

1.10) Check Replica Set

```
kubectl get rs | grep api
```

```
[root@meyclvspiapp03 ~]# kubectl get rs | grep api
firewall-api-77fbf4b5bb   1           1           1           17d
```

1.11) Describe Replica Set

```
kubectl describe rs firewall-api-77fbf4b5bb
```

```
[root@meyclvspiapp03 ~]# kubectl describe rs firewall-api-77fbf4b5bb
Name:           firewall-api-77fbf4b5bb
Namespace:      default
Selector:       app=firewall-api,pod-template-hash=77fbf4b5bb,release=stable
Labels:         app=firewall-api
                pod-template-hash=77fbf4b5bb
                release=stable
Annotations:   deployment.kubernetes.io/desired-replicas: 1
                deployment.kubernetes.io/max-replicas: 2
                deployment.kubernetes.io/revision: 1
Controlled By: Deployment/firewall-api
Replicas:       1 current / 1 desired
Pods Status:   1 Running / 0 Waiting / 0 Succeeded / 0 Failed
Pod Template:
  Labels:  app=firewall-api
            pod-template-hash=77fbf4b5bb
            release=stable
  Init Containers:
    datacontainer:
      Image:      my.registry:5000/firewall-api:v1.23.6.9
      Port:       <none>
      Host Port: <none>
      Command:
```

1.12) List all deployment in cluster

```
kubectl get deployment
kubectl get deployment | grep api
```

```
[root@meyclvspiapp03 ~]# kubectl get deployment | grep api
firewall-api          1/1     1     1           17d
```

1.13) Describe a specific deployment in cluster

```
kubectl describe deployment firewall-api
```

```
[root@myclvspiapp03 ~]# kubectl describe deployment firewall-api
Name:           firewall-api
Namespace:      default
CreationTimestamp: Thu, 24 Nov 2022 08:16:22 +0400
Labels:          <none>
Annotations:    deployment.kubernetes.io/revision: 1
Selector:        app=firewall-api,release=stable
Replicas:        1 desired | 1 updated | 1 total | 1 available | 0 unavailable
StrategyType:   RollingUpdate
MinReadySeconds: 0
RollingUpdateStrategy: 25% max unavailable, 25% max surge
Pod Template:
  Labels:  app=firewall-api
           release=stable
  Init Containers:
    datacontainer:
      Image:  my.registry:5000/firewall-api:v1.23.6.9
      Port:   <none>
      Host Port: <none>
      Command:
```

1.14) Check daemonset in cluster

```
kubectl get daemonset -n kube-system
```

```
[root@master-node1 .kube]# kubectl get daemonset -n kube-system
NAME        DESIRED   CURRENT   READY     UP-TO-DATE   AVAILABLE   NODE SELECTOR   AGE
kube-proxy   3         3         1         3            1           <none>      kubernetes.io/os=linux   20d
weave-net    3         3         0         3            0           <none>      20d
```

1.15) Describe daemonset in specific namespace

```
kubectl describe daemonset weave-net -n kube-system
```

```
[root@master-node1 ~]# kubectl describe daemonset weave-net -n kube-system
Name:           weave-net
Selector:       name=weave-net
Node-Selector: <none>
Labels:         name=weave-net
Annotations:   deprecated.daemonset.template.generation: 1
Desired Number of Nodes Scheduled: 3
Current Number of Nodes Scheduled: 3
Number of Nodes Scheduled with Up-to-date Pods: 3
Number of Nodes Scheduled with Available Pods: 0
Number of Nodes Misscheduled: 0
Pods Status:   1 Running / 2 Waiting / 0 Succeeded / 0 Failed
Pod Template:
  Labels:       name=weave-net
  Service Account: weave-net
  Init Containers:
    weave-init:
      Image:      weaveworks/weave-kube:latest
      Port:       <none>
      Host Port: <none>
      Command:
        /home/weave/init.sh
      Environment: <none>
      Mounts:
```

1.16 Check the events in a specific Namespace

```
kubectl get ev -n kube-system
```

```
[root@Master-node ~]# kubectl get ev -n kube-system
LAST SEEN TYPE      REASON          OBJECT                                MESSAGE
62s     Warning   unhealthy       pod/kube-apiserver-master-node   Liveness probe failed: Get https://192.168.0.136:6443/healthz: dial tcp 192.168.0.136:6443: connect: connection refused
62s     Normal    Pulled        pod/kube-controller-manager-master-node Container image "k8s.gcr.io/kube-controller-manager:v1.18.20" already present on machine
62s     Normal    Created       pod/kube-controller-manager-master-node Created container kube-controller-manager
62s     Normal    Started      pod/kube-controller-manager-master-node Started container kube-controller-manager
62s     Warning   Unhealthy    pod/kube-controller-manager-master-node Liveness probe failed: Get https://127.0.0.1:10257/healthz: dial tcp 127.0.0.1:10257: connect: connection refused
57: connect: connection refused
41s     Normal    LeaderElection endpoints/kube-controller-manager master-node_9a1be099-2622-45c1-a7ee-2c0399883a1 became leader
41s     Normal    LeaderElection lease/kube-controller-manager master-node_9a1be099-2622-45c1-a7ee-2c0399883a1 became leader
60s     Normal    Pulled        pod/kube-scheduler-master-node   Container image "k8s.gcr.io/kube-scheduler:v1.18.20" already present on machine
60s     Normal    Created       pod/kube-scheduler-master-node   Created container kube-scheduler
60s     Normal    Started      pod/kube-scheduler-master-node   Started container kube-scheduler
62s     Warning   Unhealthy    pod/kube-scheduler-master-node   Liveness probe failed: Get https://127.0.0.1:10259/healthz: dial tcp 127.0.0.1:10259: connect: connection refused
59: connect: connection refused
44s     Normal    LeaderElection endpoints/kube-scheduler master-node_7db21e48-1bbf-41cc-b463-58003ffeed4b became leader
44s     Normal    LeaderElection lease/kube-scheduler master-node_7db21e48-1bbf-41cc-b463-58003ffeed4b became leader
25s     Warning   InspectFailed pod/weave-net-dndk5 Failed to inspect image "weaveworks/weave-kube:latest": rpc error: code = Unknown desc = Error response from daemon: readlink /var/lib/docker/overlay2/l: invalid argument
10s     Warning   Failed       pod/weave-net-dndk5 Error: ImageInspectError
```

```
kubectl get ev -n kube-system | grep -i failed
```

```
[root@Master-node] ~# kubectl get no -n kube-system | grep -l failed
4m0s   Warning  unhealthy  pod/kube-applier-master-node
4m0s   Warning  unhealthy  pod/kube-controller-manager-master-node
4m0s   Warning  unhealthy  pod/kube-scheduler-master-node
3m0s   Warning  InspectFailed  pod/weave-net-dock
               Liveness probe failed: Get https://192.168.0.136:4643/health: dial tcp [::]:4643: i/o timeout
               Liveness probe failed: Get https://127.0.0.1:10007/ready: dial tcp [::]:10007: connect: connection refused
               Liveness probe failed: Get https://127.0.0.1:10009/health: dial tcp [::]:10009: connect: connection refused
               Failed to inspect image "weaveworks/weave-kube:latest": rpc error: code = Unknown desc = Error Response from Docker: reading
               /var/lib/docker/containers/1a2.../status: invalid argument
               Events: ImageInspectError
```

2) Cordon and UnCordon a specific Node

This command will lock the given node for creation of new Pods via scheduler . This is normally used during Maintenance

```
kubectl cordon worker-node2  
kubectl get nodes  
kubectl uncordon worker-node2
```

```
[root@master-node1 ~]# kubectl cordon worker-node2  
node/worker-node2 cordoned  
[root@master-node1 ~]# kubectl get nodes  
NAME           STATUS            ROLES      AGE   VERSION  
master-node1   NotReady        master     19d   v1.18.20  
worker-node1   NotReady        <none>    19d   v1.18.20  
worker-node2   NotReady,SchedulingDisabled  <none>    19d   v1.18.20  
[root@master-node1 ~]# kubectl uncordon worker-node2  
node/worker-node2 uncordoned  
[root@master-node1 ~]# kubectl get nodes  
NAME           STATUS    ROLES      AGE   VERSION  
master-node1   NotReady  master     19d   v1.18.20  
worker-node1   NotReady  <none>    19d   v1.18.20  
worker-node2   NotReady  <none>    19d   v1.18.20
```

3) Apply/Overwrite/Remove Label on Node

3.1) Use -L to find all nodes using disk as key , value can be any

```
kubectl label node worker-node2 disk=ssd  
kubectl get nodes -l disk=ssd  
kubectl get nodes -L disk
```

```
[root@master-node1 ~]# kubectl label node worker-node2 disk=ssd  
node/worker-node2 labeled  
[root@master-node1 ~]#  
[root@master-node1 ~]# kubectl get nodes -l disk=ssd  
NAME      STATUS    ROLES   AGE   VERSION  
worker-node2  NotReady  <none>  19d   v1.18.20  
[root@master-node1 ~]# kubectl get nodes -L disk  
NAME      STATUS    ROLES   AGE   VERSION   DISK  
master-node1  NotReady  master   19d   v1.18.20  
worker-node1  NotReady  <none>  19d   v1.18.20  
worker-node2  NotReady  <none>  19d   v1.18.20  ssd
```

3.2) Overwrite or change the value of Label on any Nodes

```
kubectl label node worker-node2 disk=sata --overwrite
```

3.3) Remove the Label disk

```
kubectl label node worker-node2 disk-
```

```
[root@master-node1 ~]# kubectl label node worker-node2 disk-  
node/worker-node2 labeled  
[root@master-node1 ~]# kubectl get nodes -L disk  
NAME      STATUS    ROLES   AGE   VERSION   DISK  
master-node1  NotReady  master   19d   v1.18.20  
worker-node1  NotReady  <none>  19d   v1.18.20  
worker-node2  NotReady  <none>  19d   v1.18.20
```

4) Create / delete Namespace

```
kubectl create ns techtracker  
kubectl get ns techtracker
```

```
[root@master-node1 .kube]# kubectl create ns techtracker
namespace/techtracker created
[root@master-node1 .kube]# kubectl get ns techtracker
NAME      STATUS   AGE
techtracker   Active   14s
```

5) Manage cluster from Local host using kubeconfig

```
mkdir -p $HOME/.kube
cd $HOME/.kube/
vi config
copy paste the content of /etc/kubernetes/admin.conf and save
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

6) Create / delete / Rollout / Rollback Deployment

6.1) Prepare Deployment YAML

```
vi deployment.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: techtracker-deployment
  namespace: techtracker
  labels:
    app: nginx
    env: qa
spec:
  replicas: 10
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.14.2
          ports:
            - containerPort: 80
```

6.2) Open 2nd vertical terminal and execute live watch to see pods getting created

```
# On Terminal 2  
watch -n 1 kubectl get all -n techtracker
```

6.3) Create Deployment using Apply :

```
# On Terminal 1  
kubectl apply -f deployment.yaml
```

```
[root@master-node yaml]# kubectl apply -f deployment.yaml --record  
deployment.apps/techtracker-deployment created  
[root@master-node yaml]#  
  
Every 1.0s: kubectl get all -n techtracker Sun Dec 11 07:26:38 2022  
NAME READY STATUS RESTARTS AGE  
pod/techtracker-deployment-6b474476c4-2954l 1/1 Running 0 4s  
pod/techtracker-deployment-6b474476c4-46rct 1/1 Running 0 4s  
pod/techtracker-deployment-6b474476c4-6qpk5 1/1 ContainerCreating 0 4s  
pod/techtracker-deployment-6b474476c4-8vqq5 1/1 ContainerCreating 0 4s  
pod/techtracker-deployment-6b474476c4-9p4ls 1/1 ContainerCreating 0 4s  
pod/techtracker-deployment-6b474476c4-k2fzl 1/1 ContainerCreating 0 4s  
pod/techtracker-deployment-6b474476c4-k9cjw 1/1 ContainerCreating 0 4s  
pod/techtracker-deployment-6b474476c4-lrjsz 1/1 ContainerCreating 0 4s  
pod/techtracker-deployment-6b474476c4-mw9kd 1/1 ContainerCreating 0 4s  
pod/techtracker-deployment-6b474476c4-vzxc4 1/1 ContainerCreating 0 4s  
  
NAME READY UP-TO-DATE AVAILABLE AGE  
deployment.apps/techtracker-deployment 5/10 10 5 4s  
  
NAME DESIRED CURRENT READY AGE  
replicaset.apps/techtracker-deployment-6b474476c4 10 10 10 4s
```

Every 1.0s: kubectl get all -n techtracker Sun Dec 11 07:26:08 2022

NAME	READY	STATUS	RESTARTS	AGE
pod/techtracker-deployment-6b474476c4-2954l	1/1	Running	0	34s
pod/techtracker-deployment-6b474476c4-46rct	1/1	Running	0	34s
pod/techtracker-deployment-6b474476c4-6qpk5	1/1	Running	0	34s
pod/techtracker-deployment-6b474476c4-8vqq5	1/1	Running	0	34s
pod/techtracker-deployment-6b474476c4-9p4ls	1/1	Running	0	34s
pod/techtracker-deployment-6b474476c4-k2fzl	1/1	Running	0	34s
pod/techtracker-deployment-6b474476c4-k9cjw	1/1	Running	0	34s
pod/techtracker-deployment-6b474476c4-lrjsz	1/1	Running	0	34s
pod/techtracker-deployment-6b474476c4-mw9kd	1/1	Running	0	34s
pod/techtracker-deployment-6b474476c4-vzxc4	1/1	Running	0	34s

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
deployment.apps/techtracker-deployment	10/10	10	10	34s

NAME	DESIRED	CURRENT	READY	AGE
replicaset.apps/techtracker-deployment-6b474476c4	10	10	10	34s

```
kubectl get deployment techtracker-deployment -n techtracker
```

```
[root@master-node1 yaml]# kubectl get deployment techtracker-deployment -n techtracker
NAME           READY   UP-TO-DATE   AVAILABLE   AGE
techtracker-deployment   10/10    10          10         107s
```

6.4) Check the Container Image version is correct :

```
kubectl describe pod techtracker-deployment-6b474476c4-8vqq5 -n techtracker | grep -i image
```

```
[root@master-node1 yaml]# kubectl describe pod techtracker-deployment-6b474476c4-8vqq5 -n techtracker | grep -i image
  Image:      nginx:1.14.2
  Image ID:   docker-pullable://nginx@sha256:f7908fb6c02e0ce69257d9bd9cf37ae20a60f1df7563c3aza6abe24160306b8d
  Normal  Pulled  5m48s  kubelet        Container image "nginx:1.14.2" already present on machine
```

6.5) Check the revision no. in rollout history of deployment

```
kubectl rollout history deployment techtracker-deployment -n techtracker
```

```
[root@master-node1 ~]# kubectl rollout history deployment techtracker-deployment -n techtracker
deployment.apps/techtracker-deployment
REVISION  CHANGE-CAUSE
1        kubectl apply --filename=deployment.yaml --record=true
```

6.6) Check the Rolling Update Strategy of deployment

```
spec:
  minReadySeconds: 20
  progressDeadlineSeconds: 600
  replicas: 4
  revisionHistoryLimit: 10
  selector:
    matchLabels:
      name: webapp
  strategy:
    rollingUpdate:
      maxSurge: 25%
      maxUnavailable: 25%
    type: RollingUpdate
```

```
kubectl describe deployment techtracker-deployment -n techtracker | grep -i RollingUpdateStrategy
```

```
[root@master-node1 ~]# kubectl describe deployment techtracker-deployment -n techtracker | grep -i RollingUpdateStrategy
RollingUpdateStrategy: 25% max unavailable, 25% max surge
```

6.7) Set latest image in deployment rollout

```
kubectl set image deployment techtracker-deployment nginx=nginx:latest -n techtracker --record
```

```
[root@master-node1 ~]# kubectl set image deployment techtracker-deployment nginx=nginx:latest -n techtracker --record
deployment.apps/techtracker-deployment image updated
[root@master-node1 ~]# Every 1.0s: kubectl get all -n techtracker
NAME                                         READY   STATUS    RESTARTS   AGE
pod/techtracker-deployment-684bd4f4c-59bwn   0/1     Pending   0          35s
pod/techtracker-deployment-684bd4f4c-c7n8q   0/1     Pending   0          38s
pod/techtracker-deployment-684bd4f4c-g2ssv   0/1     Pending   0          28s
pod/techtracker-deployment-684bd4f4c-gltxl   0/1     Pending   0          34s
pod/techtracker-deployment-684bd4f4c-s4r55   0/1     Pending   0          34s
pod/techtracker-deployment-6b474476c4-2954l  1/1     Running  0          22m
pod/techtracker-deployment-6b474476c4-46rct  1/1     Running  0          22m
pod/techtracker-deployment-6b474476c4-6qpk5  1/1     Terminating 0          22m
pod/techtracker-deployment-6b474476c4-8vqq5  1/1     Running  0          22m
pod/techtracker-deployment-6b474476c4-k2fzl  1/1     Running  0          22m
pod/techtracker-deployment-6b474476c4-k9cjw  1/1     Running  0          22m
pod/techtracker-deployment-6b474476c4-lrjsz   1/1     Running  0          22m
pod/techtracker-deployment-6b474476c4-mw9kd  1/1     Running  0          22m
pod/techtracker-deployment-6b474476c4-vzxc4  1/1     Running  0          22m
                                                 STATUS    UP-TO-DATE   AVAILABLE   AGE
deployment.apps/techtracker-deployment        0/10    5           0          22m
                                                 DESIRED   CURRENT    READY      AGE
replicaset.apps/techtracker-deployment-684bd4f4c  5         5         0          36s
replicaset.apps/techtracker-deployment-6b474476c4  8         8         0          22m
NAME                                         READY   STATUS    RESTARTS   AGE
pod/techtracker-deployment-684bd4f4c-59bwn   0/1     Pending   0          6m22s
pod/techtracker-deployment-684bd4f4c-c7n8q   0/1     Pending   0          6m22s
pod/techtracker-deployment-684bd4f4c-g2ssv   0/1     Pending   0          6m21s
pod/techtracker-deployment-684bd4f4c-gltxl   0/1     Pending   0          6m22s
pod/techtracker-deployment-684bd4f4c-s4r55   0/1     Pending   0          6m21s
pod/techtracker-deployment-6b474476c4-2954l  1/1     Running  0          22m
pod/techtracker-deployment-6b474476c4-46rct  1/1     Running  0          22m
pod/techtracker-deployment-6b474476c4-6qpk5  1/1     Terminating 0          22m
pod/techtracker-deployment-6b474476c4-8vqq5  1/1     Running  0          22m
pod/techtracker-deployment-6b474476c4-9p4ls  1/1     Terminating 0          22m
pod/techtracker-deployment-6b474476c4-k2fzl  1/1     Running  0          22m
pod/techtracker-deployment-6b474476c4-k9cjw  1/1     Running  0          22m
pod/techtracker-deployment-6b474476c4-lrjsz   1/1     Running  0          22m
pod/techtracker-deployment-6b474476c4-mw9kd  1/1     Running  0          22m
pod/techtracker-deployment-6b474476c4-vzxc4  1/1     Running  0          22m
                                                 STATUS    UP-TO-DATE   AVAILABLE   AGE
deployment.apps/techtracker-deployment        0/10    5           0          22m
                                                 DESIRED   CURRENT    READY      AGE
replicaset.apps/techtracker-deployment-684bd4f4c  5         5         0          6m22s
replicaset.apps/techtracker-deployment-6b474476c4  8         8         0          22m
```

Every 1.0s: kubectl get all -n techtracker				
NAME	READY	STATUS	RESTARTS	AGE
pod/techtracker-deployment-684bd4f4c-59bwn	1/1	Running	0	14m
pod/techtracker-deployment-684bd4f4c-6f2gp	1/1	Running	0	2m59s
pod/techtracker-deployment-684bd4f4c-79z4b	1/1	Running	0	3m1s
pod/techtracker-deployment-684bd4f4c-c7n8q	1/1	Running	0	14m
pod/techtracker-deployment-684bd4f4c-g2ssv	1/1	Running	0	14m
pod/techtracker-deployment-684bd4f4c-gltxl	1/1	Running	0	14m
pod/techtracker-deployment-684bd4f4c-kql9v	1/1	Running	0	3m3s
pod/techtracker-deployment-684bd4f4c-ngqfx	1/1	Running	0	2m59s
pod/techtracker-deployment-684bd4f4c-s4r55	1/1	Running	0	14m
pod/techtracker-deployment-684bd4f4c-vjl4g	1/1	Running	0	3m1s

6.8) Check the 2nd Revision added in rollout history:

```
kubectl rollout history deployment techtracker-deployment -n techtracker
```

```
[root@master-node1 ~]# kubectl rollout history deployment techtracker-deployment -n techtracker
deployment.apps/techtracker-deployment
REVISION  CHANGE-CAUSE
1          kubectl apply --filename=deployment.yaml --record=true
2          kubectl set image deployment techtracker-deployment nginx=nginx:latest --namespace=techtracker --record=true
```

6.9) Check/Verify the Container Image should be Latest :

```
kubectl describe pod techtracker-deployment-6b474476c4-8vqq5 -n techtracker | grep -i image
```

```
[root@master-node1 ~]# kubectl describe pod techtracker-deployment-684bd4f4c-59bwn -n techtracker | grep -i image
  Image:      nginx:latest
  Image ID:   docker-pullable://nginx@sha256:ab589a3c466e347b1c0573be23356676df90cd7ce2dbf6ec332a5f0a8b5e59db
Normal    Pulling      2m16s      kubelet      Pulling image "nginx:latest"
Normal    Pulled       119s      kubelet      Successfully pulled image "nginx:latest"
```

6.10) Rollback to a specific old version using -to-revision=1

```
kubectl rollout undo deployment techtracker-deployment -n techtracker --to-revision=1
```

```
[root@master-node1 ~]# kubectl rollout history deployment techtracker-deployment -n techtracker
deployment.apps/techtracker-deployment
REVISION: CHANGE-CASE
1  kubectl apply --filename=deployment.yaml --record=true
2  kubectl set image deployment techtracker-deployment nginx:nginx:latest --name=techtracker --record=true
[root@master-node1 ~]# kubectl rollout undo deployment techtracker-deployment -n techtracker --to-revision=1
deployment.apps/techtracker-deployment rolled back
[root@master-node1 ~]# Every 1.0s: kubectl get all -n techtracker
NAME                                     READY   STATUS    RESTARTS   AGE
pod/techtracker-deployment-684dd14c-59bw   1/1    Running   0          10m
pod/techtracker-deployment-684dd14c-672gp  1/1    Terminating   0          6m42s
...
pod/techtracker-deployment-684dd14c-79z4   1/1    Running   0          6m42s
pod/techtracker-deployment-684dd14c-c7nbq  0/1    Terminating   0          10m
pod/techtracker-deployment-684dd14c-q25sv  0/1    Terminating   0          10m
pod/techtracker-deployment-684dd14c-q1txl  1/1    Running   0          10m
pod/techtracker-deployment-684dd14c-kql9x  1/1    Running   0          6m44s
...
pod/techtracker-deployment-684dd14c-ngxts  1/1    Running   0          6m40s
pod/techtracker-deployment-684dd14c-skr55  1/1    Running   0          10m
pod/techtracker-deployment-684dd14c-vjl8g  1/1    Running   0          6m42s
...
pod/techtracker-deployment-6b474476c4-29541 1/1    Terminating   0          34m
pod/techtracker-deployment-6b474476c4-46rct 1/1    Terminating   0          34m
pod/techtracker-deployment-6b474476c4-48xkp 0/1    ContainerCreating   0          2s
pod/techtracker-deployment-6b474476c4-8vgpb 1/1    Terminating   0          34m
pod/techtracker-deployment-6b474476c4-4t6x  0/1    ContainerCreating   0          2s
pod/techtracker-deployment-6b474476c4-avjvx 0/1    ContainerCreating   0          0s
pod/techtracker-deployment-6b474476c4-aw8kd 1/1    Terminating   0          34m
pod/techtracker-deployment-6b474476c4-6yfrd 0/1    ContainerCreating   0          2s
pod/techtracker-deployment-6b474476c4-qgr9c 0/1    ContainerCreating   0          2s
pod/techtracker-deployment-6b474476c4-vqz05 1/1    Running   0          2s
pod/techtracker-deployment-6b474476c4-vzx4c 1/1    Terminating   0          34m
NAME                                     READY   UP-TO-DATE   AVAILABLE   AGE
deployment.apps/techtracker-deployment   8/8h   6          8          34m
NAME                                     DESIRED   CURRENT   READY   AGE
replicaset.apps/techtracker-deployment-684dd14c 7        7        7        10m
replicaset.apps/techtracker-deployment-6b474476c4 6        6        3        34m
```

6.11) Check and verify the container version after rollback

```
kubectl describe pod techtracker-deployment-6b474476c4-dld8w -n techtracker | grep -i image
```

```
[root@master-node1 ~]# kubectl describe pod techtracker-deployment-6b474476c4-dld8w -n techtracker | grep -i image
  Image:           nginx:1.14.2
  Image ID:       docker-pullable://nginx@sha256:f7988fb6c02e0ce69257d9bd9cf37ae20a60f1df7563c3a2a6abe24160306b8d
  Normal   Pulled      12m   kubelet   Container image "nginx:1.14.2" already present on machine
[root@master-node1 ~]#
```

6.12) Delete a existing deployment

```
kubectl delete deployment techtracker-deployment -n techtracker
```

```
[root@master-node1 ~]# kubectl delete deployment techtracker-deployment -n techtracker
deployment.apps "techtracker-deployment" deleted
[root@master-node1 ~]# Every 1.0s: kubectl get all -n techtracker
NAME                                     READY   STATUS    RESTARTS   AGE
pod/techtracker-deployment-6b474476c4-29541 1/1    Terminating   0          49m
pod/techtracker-deployment-6b474476c4-46rct 1/1    Terminating   0          49m
pod/techtracker-deployment-6b474476c4-48xkp 1/1    Terminating   0          15m
pod/techtracker-deployment-6b474476c4-8vgpb 1/1    Terminating   0          49m
pod/techtracker-deployment-6b474476c4-4t6x  1/1    Terminating   0          15m
pod/techtracker-deployment-6b474476c4-aw8kd 1/1    Terminating   0          15m
pod/techtracker-deployment-6b474476c4-6yfrd 1/1    Terminating   0          15m
pod/techtracker-deployment-6b474476c4-7an22 1/1    Terminating   0          15m
pod/techtracker-deployment-6b474476c4-7dkx  1/1    Terminating   0          15m
pod/techtracker-deployment-6b474476c4-8rjxx 1/1    Terminating   0          15m
pod/techtracker-deployment-6b474476c4-9p9kd 1/1    Terminating   0          49m
pod/techtracker-deployment-6b474476c4-9pjh1b 1/1    Terminating   0          15m
pod/techtracker-deployment-6b474476c4-s6frd 1/1    Terminating   0          15m
pod/techtracker-deployment-6b474476c4-seqic 1/1    Terminating   0          15m
pod/techtracker-deployment-6b474476c4-vgr65 1/1    Terminating   0          15m
pod/techtracker-deployment-6b474476c4-vzx4c 1/1    Terminating   0          49m
```

7) Delete the Stuck pods in Terminating State

```
[root@master-node1 ~]# kubectl get all -n techtracker
NAME                               READY   STATUS        RESTARTS   AGE
pod/techtracker-deployment-6b474476c4-46rct   1/1    Terminating   0          74m
pod/techtracker-deployment-6b474476c4-4mxkq    1/1    Terminating   0          39m
pod/techtracker-deployment-6b474476c4-8vqq5    1/1    Terminating   0          74m
pod/techtracker-deployment-6b474476c4-d48fb   1/1    Terminating   0          39m
pod/techtracker-deployment-6b474476c4-dld8w   1/1    Terminating   0          39m
pod/techtracker-deployment-6b474476c4-fmn22   1/1    Terminating   0          39m
pod/techtracker-deployment-6b474476c4-fx9cx   1/1    Terminating   0          39m
pod/techtracker-deployment-6b474476c4-mvjvx   1/1    Terminating   0          39m
pod/techtracker-deployment-6b474476c4-mw9kd   1/1    Terminating   0          74m
pod/techtracker-deployment-6b474476c4-nbjtb   1/1    Terminating   0          39m
pod/techtracker-deployment-6b474476c4-s6frd   1/1    Terminating   0          39m
pod/techtracker-deployment-6b474476c4-sqr9c   1/1    Terminating   0          39m
pod/techtracker-deployment-6b474476c4-vqz65   1/1    Terminating   0          39m
pod/techtracker-deployment-6b474476c4-vzxc4   1/1    Terminating   0          74m
[root@master-node1 ~]#
```

```
kubectl delete --grace-period=0 --force --namespace techtracker $(kubectl get pods -n techtracker -o name)
```

```
[root@master-node1 ~]# kubectl delete --grace-period=0 --force --namespace techtracker $(kubectl get pods -n techtracker -o name)
warning: Immediate deletion does not wait for confirmation that the running resource has been terminated. The resource may continue to run on the cluster indefinitely.
pod "TechTracker-deployment-6b474476c4-46rct" force deleted
pod "TechTracker-deployment-6b474476c4-4mxkq" force deleted
pod "TechTracker-deployment-6b474476c4-8vqq5" force deleted
pod "TechTracker-deployment-6b474476c4-d48fb" force deleted
pod "TechTracker-deployment-6b474476c4-dld8w" force deleted
pod "TechTracker-deployment-6b474476c4-fmn22" force deleted
pod "TechTracker-deployment-6b474476c4-fx9cx" force deleted
pod "TechTracker-deployment-6b474476c4-mvjvx" force deleted
pod "TechTracker-deployment-6b474476c4-mw9kd" force deleted
pod "TechTracker-deployment-6b474476c4-nbjtb" force deleted
pod "TechTracker-deployment-6b474476c4-s6frd" force deleted
pod "TechTracker-deployment-6b474476c4-sqr9c" force deleted
pod "TechTracker-deployment-6b474476c4-vqz65" force deleted
pod "TechTracker-deployment-6b474476c4-vzxc4" force deleted
```

8) Edit deployment to change any values like maxSurge or maxUnavailable

```
kubectl edit deployment techtracker-deployment -n techtracker
```

```
strategy:
  rollingUpdate:
    maxSurge: 40%
    maxUnavailable: 10%
  type: RollingUpdate
template:
  metadata:
    creationTimestamp: null
```

9) Create/Delete/Verify pods using DaemonSet apply

9.1) Prepare DaemonSet yaml .

```
vi daemonset.yaml
apiVersion: apps/v1
kind: DaemonSet
metadata:
  name: techtracker-daemonset
  namespace: techtracker
spec:
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.14.2
          ports:
            - containerPort: 80
```

9.2) Execute DaemonSet yaml using apply

```
kubectl apply -f daemonset.yaml
watch -n 1 kubectl get all -n techtracker -o wide
```

nginx pod created on each worker node.

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
pod/techtracker-daemonset-5tnwg	1/1	Running	0	74s	10.44.0.1	worker-node2	<none>	<none>
pod/techtracker-daemonset-8t9k2	1/1	Running	0	74s	10.32.0.6	worker-node1	<none>	<none>
NAME	DESIRED	CURRENT	READY	UP-TO-DATE	AVAILABLE	NODE SELECTOR	AGE	CONTAINERS IMAGES SELECTOR
daemonset.apps/techtracker-daemonset	2	2	2	2	2	<none>	74s	nginx nginx:1.14.2 app=nginx

9.3) Check DaemonSet

```
kubectl get daemonset -n techtracker
```

```
[root@master-node1 ~]# kubectl get daemonset -n techtracker
NAME           DESIRED   CURRENT   READY   UP-TO-DATE   AVAILABLE   NODE SELECTOR   AGE
techtracker-daemonset   2         2         2        2           2           <none>      9m
```

9.4) Describe DaemonSet

```
kubectl describe daemonset techtracker-daemonset -n techtracker
```

```
[root@master-node1 ~]# kubectl describe daemonset techtracker-daemonset -n techtracker
Name:           techtracker-daemonset
Selector:       app=nginx
Node-Selector:  <none>
Labels:         <none>
Annotations:   deprecated.daemonset.template.generation: 1
Desired Number of Nodes Scheduled: 1
Current Number of Nodes Scheduled: 1
Number of Nodes Scheduled with Up-to-date Pods: 1
Number of Nodes Scheduled with Available Pods: 1
Number of Nodes Misscheduled: 1
Pods Status:   2 Running / 0 Waiting / 0 Succeeded / 0 Failed
Pod Template:
  Labels:  app=nginx
  Containers:
    nginx:
      Image:      nginx:1.14.2
      Port:       80/TCP
      Host Port:  0/TCP
      Environment: <none>
      Mounts:     <none>
      Volumes:    <none>
  Events:
    Type  Reason          Age   From            Message
    ----  ----          --   --              --
    Normal SuccessfulCreate 12m  daemonset-controller  Created pod: techtracker-daemonset-8t9k2
    Normal SuccessfulCreate 12m  daemonset-controller  Created pod: techtracker-daemonset-5tnwg
[root@master-node1 ~]#
```

9.5) Delete DaemonSet

```
kubectl delete daemonset techtracker-daemonset -n techtracker
```

```
[root@master-node1 ~]# kubectl delete daemonset techtracker-daemonset -n techtracker
daemonset.apps "techtracker-daemonset" deleted
[root@master-node1 ~]# kubectl get all -n techtracker -o wide
NAME                   READY   STATUS    RESTARTS   AGE   IP           NODE   NOMINATED NODE   READINESS GATES
pod/techtracker-daemonset-8t9k2  1/1    Terminating   0    19m  10.32.0.6  worker-node1  <none>  <none>
```

9.6) Nodeselector Scheduling

```
kubectl label node worker-node1 env=test  
kubectl label node worker-node2 env=prod
```

```
[root@master-node1 ~]# kubectl get nodes  
NAME        STATUS   ROLES    AGE     VERSION  
master-node1  NotReady master   23d    v1.18.20  
worker-node1  NotReady <none>   23d    v1.18.20  
worker-node2  NotReady <none>   23d    v1.18.20  
[root@master-node1 ~]# kubectl label node worker-node1 env=test  
node/worker-node1 labeled  
[root@master-node1 ~]# kubectl label node worker-node2 env=prod  
node/worker-node2 labeled
```

```
apiVersion: apps/v1  
kind: Deployment  
metadata:  
  name: techtracker-deployment  
  namespace: techtracker  
spec:  
  replicas: 6  
  selector:  
    matchLabels:  
      app: nginx  
  template:  
    metadata:  
      labels:  
        app: nginx  
    spec:  
      containers:  
      - name: nginx  
        image: nginx:1.14.2  
        ports:  
        - containerPort: 80  
      nodeSelector:  
        env: test  
  kubectl apply -f schedulor_deployment.yaml
```

All pods created on worker-node1 because we had set nodeSelector env:test and env:test label on worker-node1

Every 1.0s: kubectl get all -n techtracker -o wide											Wed Dec 14 11:45:10 -0500 2019
NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS	GATES		
pod/techtracker-deployment-98f95458-47fkf	0/1	ContainerCreating	0	8s	<none>	worker-node1	<none>	<none>	<none>		
pod/techtracker-deployment-98f95458-6vv26	0/1	ContainerCreating	0	8s	<none>	worker-node1	<none>	<none>	<none>		
pod/techtracker-deployment-98f95458-7lkj7	0/1	ContainerCreating	0	8s	<none>	worker-node1	<none>	<none>	<none>		
pod/techtracker-deployment-98f95458-mzrg	0/1	ContainerCreating	0	8s	<none>	worker-node1	<none>	<none>	<none>		
pod/techtracker-deployment-98f95458-pjppg	0/1	ContainerCreating	0	8s	<none>	worker-node1	<none>	<none>	<none>		
pod/techtracker-deployment-98f95458-scqjg	0/1	ContainerCreating	0	8s	<none>	worker-node1	<none>	<none>	<none>		
NAME	READY	UP-TO-DATE	AVAILABLE	AGE	CONTAINERS	IMAGES	SELECTOR				
deployment.apps/techtracker-deployment	0/6	6	0	8s	nginx	nginx:1.14.2	app=nginx				
NAME	DESIRED	CURRENT	READY	AGE	CONTAINERS	IMAGES	SELECTOR				
replicaset.apps/techtracker-deployment-98f95458	6	6	0	8s	nginx	nginx:1.14.2	app=nginx,pod-template-hash=98f95458				

9.7) nodeName Scheduling

```

apiVersion: apps/v1
kind: Deployment
metadata:
  name: techtracker-deployment
  namespace: techtracker
spec:
  replicas: 6
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.14.2
          ports:
            - containerPort: 80
  nodeName: worker-node2
  kubectl apply -f scheduler_deployment.yaml

```

All pods will get created on worker-node2

Every 1.0s: kubectl get all -n techtracker -o wide											Wed Dec 14 11:45:10 -0500 2019
NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS	GATES		
pod/techtracker-deployment-7fcc6f7d5-8nkvl	1/1	Running	0	11s	10.44.0.6	worker-node2	<none>	<none>	<none>		
pod/techtracker-deployment-7fcc6f7d5-dbwmh	1/1	Running	0	11s	10.44.0.2	worker-node2	<none>	<none>	<none>		
pod/techtracker-deployment-7fcc6f7d5-hptrd	1/1	Running	0	11s	10.44.0.4	worker-node2	<none>	<none>	<none>		
pod/techtracker-deployment-7fcc6f7d5-p4h8s	1/1	Running	0	11s	10.44.0.5	worker-node2	<none>	<none>	<none>		
pod/techtracker-deployment-7fcc6f7d5-prwdx	1/1	Running	0	11s	10.44.0.3	worker-node2	<none>	<none>	<none>		
pod/techtracker-deployment-7fcc6f7d5-s2npw	1/1	Running	0	11s	10.44.0.1	worker-node2	<none>	<none>	<none>		
NAME	READY	UP-TO-DATE	AVAILABLE	AGE	CONTAINERS	IMAGES	SELECTOR				
deployment.apps/techtracker-deployment	6/6	6	6	11s	nginx	nginx:1.14.2	app=nginx				
NAME	DESIRED	CURRENT	READY	AGE	CONTAINERS	IMAGES	SELECTOR				
replicaset.apps/techtracker-deployment-7fcc6f7d5	6	6	6	11s	nginx	nginx:1.14.2	app=nginx,pod-template-hash=7fcc6f7d5				

10) Taint and Toleration

By Default Scheduler never created pods on tainted node . But if we still want some heavy or specific application to be deployed on tainted node then we can use toleration in deployment file for that application .

- Cordon cannot be ByPassed
- Taint can be ByPassed using Toleration

Check the tainted Nodes in Cluster

```
kubectl describe no master-node1 -n techtracker | grep -i taint
```

```
[root@master-node1 ~]# kubectl describe no master-node1 -n techtracker | grep -i taint
Taints:           node.kubernetes.io/not-ready:NoExecute
```

taint of Effect — NoExecute is applied at master node

Apply taint on node1

```
kubectl taint node worker-node1 dubai=hot:NoSchedule
```

```
[root@master-node1 ~]# kubectl taint node worker-node1 dubai=hot:NoSchedule
node/worker-node1 tainted
[root@master-node1 ~]#
```

```
kubectl describe no worker-node1 -n techtracker | grep -i taint
```

```
[root@master-node1 ~]# kubectl describe no worker-node1 -n techtracker | grep -i taint
Taints:          node.kubernetes.io/unreachable:NoExecute
```

#Remove taint from node

```
kubectl taint node worker-node1 node.kubernetes.io/unreachable:NoExecute-
```

```
[root@master-node1 ~]# kubectl taint node worker-node1 node.kubernetes.io/unreachable:NoExecute-
node/worker-node1 untainted
```

Create nodes on tainted node

```
[root@master-node1 yaml]# kubectl describe nodes worker-node1 -n techtracker
Name:           worker-node1
Roles:          <none>
Labels:         beta.kubernetes.io/arch=amd64
                beta.kubernetes.io/os=linux
                env=test
                kubernetes.io/arch=amd64
                kubernetes.io/hostname=worker-node1
                kubernetes.io/os=linux
Annotations:   kubeadm.alpha.kubernetes.io/cri-socket: /var/run/dockershim.sock
                node.alpha.kubernetes.io/ttl: 0
                volumes.kubernetes.io/controller-managed-attach-detach: true
CreationTimestamp: Sun, 20 Nov 2022 14:43:35 -0500
Taints:         dubai=hot:NoSchedule
Unschedulable:  false
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: techtracker-deployment
  namespace: techtracker
spec:
  replicas: 6
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
```

```
containers:
- name: nginx
  image: nginx:1.14.2
  ports:
  - containerPort: 80
tolerations:
- key: "dubai"
  operator: "Equal"
  value: "hot"
  effect: "NoSchedule"
~  
kubectl apply -f taint_toleration.yaml
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE
pod/techtracker-deployment-79f8f98c5f-9h855	1/1	Running	0	19s	10.32.0.2	worker-node1
pod/techtracker-deployment-79f8f98c5f-dd5nl	1/1	Running	0	19s	10.32.0.4	worker-node1
pod/techtracker-deployment-79f8f98c5f-dstbl	1/1	Running	0	19s	10.32.0.6	worker-node1
pod/techtracker-deployment-79f8f98c5f-dtbbx	1/1	Running	0	19s	10.32.0.7	worker-node1
pod/techtracker-deployment-79f8f98c5f-mnfhp	1/1	Running	0	19s	10.32.0.5	worker-node1
pod/techtracker-deployment-79f8f98c5f-vsn9v	1/1	Running	0	19s	10.32.0.3	worker-node1

=====

=====

11. Affinity and Antiaffinity

Apply label app=store on which customer want to deploy redis pod

```
kubectl label node worker-node1 app=store  
kubectl label node worker-node2 app=store  
kubectl get nodes -L app
```

```
[root@master-node1 yaml]# kubectl label node worker-node1 app=store
node/worker-node1 labeled
[root@master-node1 yaml]# kubectl get nodes -L app
NAME      STATUS   ROLES   AGE    VERSION   APP
master-node1  NotReady  master  28d    v1.18.20
worker-node1  Ready     <none>  28d    v1.18.20  store
worker-node2  Ready     <none>  28d    v1.18.20
[root@master-node1 yaml]# kubectl label node worker-node2 app=store
node/worker-node2 labeled
[root@master-node1 yaml]# kubectl get nodes -L app
NAME      STATUS   ROLES   AGE    VERSION   APP
master-node1  NotReady  master  28d    v1.18.20
worker-node1  Ready     <none>  28d    v1.18.20  store
worker-node2  Ready     <none>  28d    v1.18.20  store
[root@master-node1 yaml]# █
```

Apply Antiaffinity yaml and check expected output

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: redis-cache
  namespace: techtracker
spec:
  selector:
    matchLabels:
      app: store
  replicas: 3
  template:
    metadata:
      labels:
        app: store
    spec:
      affinity:
        podAntiAffinity:
          requiredDuringSchedulingIgnoredDuringExecution:
            - labelSelector:
                matchExpressions:
                  - key: app
                    operator: In
                    values:
                      - store
              topologyKey: "kubernetes.io/hostname"
  containers:
    - name: redis-server
      image: redis:3.2-alpine
```

```
kubectl apply -f antiaffinity.yaml  
kubectl get all -n techtracker
```

```
[root@master-node1 yaml]# kubectl get all -n techtracker -o wide  
NAME          READY   STATUS    RESTARTS   AGE     IP           NODE      NOMINATED NODE   READINESS GATES  
pod/redis-cache-7474cd8cdd-p5lpk  1/1    Running   0          4m7s   10.44.0.2    worker-node2 <none>        <none>  
pod/redis-cache-7474cd8cdd-qhrtw  1/1    Running   0          4m7s   10.32.0.2    worker-node1 <none>        <none>  
pod/redis-cache-7474cd8cdd-zs7c8  0/1    Pending    0          4m7s   <none>       <none>      <none>  
  
NAME          READY   UP-TO-DATE  AVAILABLE  AGE     CONTAINERS   IMAGES          SELECTOR  
deployment.apps/redis-cache     2/3     2          2          4m7s   redis-server  redis:3.2-alpine  app=store  
  
NAME          DESIRED  CURRENT  READY     AGE     CONTAINERS   IMAGES          SELECTOR  
replicaset.apps/redis-cache-7474cd8cdd  3        3         2          4m7s   redis-server  redis:3.2-alpine  app=store,pod-template-hash=7474cd8cdd
```

Its showing 2/3 and 1 pod is pending because there are only 2 nodes having label app=store and scheduler will not create 2 similar pods on 1 worker node .

Apply Affinity and Antiaffinity yaml and check output

```
apiVersion: apps/v1  
kind: Deployment  
metadata:  
  name: web-server  
spec:  
  selector:  
    matchLabels:  
      app: web-store  
  replicas: 3  
  template:  
    metadata:  
      labels:  
        app: web-store  
    spec:  
      affinity:  
        podAntiAffinity:  
          requiredDuringSchedulingIgnoredDuringExecution:  
          - labelSelector:  
              matchExpressions:  
              - key: app  
                operator: In  
                values:  
                - web-store  
              topologyKey: "kubernetes.io/hostname"  
        podAffinity:  
          requiredDuringSchedulingIgnoredDuringExecution:  
          - labelSelector:  
              matchExpressions:  
              - key: app  
                operator: In
```

```
values:
  - store
  topologyKey: "kubernetes.io/hostname"
containers:
  - name: web-app
    image: nginx:1.16-alpine
kubectl apply -f affinityantiaffinity.yaml
kubectl get all -n techtracker -o wide
```

```
[root@master-node1 yaml]# kubectl get all -n techtracker -o wide
NAME                                         READY   STATUS    RESTARTS   AGE   IP           NODE   NOMINATED-NODE   READINESS   GATES
pod/redis-cache-7474cd8cdd-p5lpk   1/1    Running   0          16m   10.44.0.2   worker-node2   <none>        <none>
pod/redis-cache-7474cd8cdd-qhrtw  1/1    Running   0          16m   10.32.0.2   worker-node1   <none>        <none>
pod/redis-cache-7474cd8cdd-zs7c8  0/1    Pending   0          16m   <none>       <none>        <none>
pod/web-server-7f49979444-lbwir  1/1    Running   0          2m2s  10.32.0.3   worker-node1   <none>        <none>
pod/web-server-7f49979444-ghs86  1/1    Running   0          2m2s  10.44.0.5   worker-node2   <none>        <none>

NAME                           READY   UP-TO-DATE   AVAILABLE   AGE   CONTAINERS   IMAGES          SELECTOR
deployment.apps/redis-cache  2/3     3            2           16m   redis-server   redis:3.2-alpine   app=store
deployment.apps/web-server   2/2     2            2           2m2s  web-app      nginx:1.16-alpine   app=web-store

NAME                           DESIRED   CURRENT   READY   AGE   CONTAINERS   IMAGES          SELECTOR
replicaset.apps/redis-cache  3         3           2           16m   redis-server   redis:3.2-alpine   app=store,pod-template-hash=7474cd8cdd
replicaset.apps/web-server   2         2           2           2m2s  web-app      nginx:1.16-alpine   app=web-store,pod-template-hash=7f49979444
```

web-server pod deployed on node where redis is already deployed

```
===== ===== ===== ===== ===== ===== =====
```

12. Secrets and Config Maps

Deployment Yaml for creating a mysql pod

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: mysql
  namespace: techtracker
spec:
  selector:
    matchLabels:
      app: mysql
  template:
    metadata:
      labels:
        app: mysql
    spec:
      containers:
        - image: mysql:5.6
```

```
name: mysql
env:
  - name: MYSQL_ROOT_PASSWORD
    value: password
kubectl apply -f resource_secret.yaml
watch -n 1 kubectl get all -n techtracker
```

NAME	READY	STATUS	RESTARTS	AGE
pod/mysql-7bc745d4f8-wnr2p	1/1	Running	0	4m34s
NAME	READY	UP-TO-DATE	AVAILABLE	AGE
deployment.apps/mysql	1/1	1	1	4m34s
NAME	DESIRED	CURRENT	READY	AGE
replicaset.apps/mysql-7bc745d4f8	1	1	1	4m34s

env variable value password is visible , anyone who has namespace access can view the secret values which is not recommended. That's why we have to use secret and store values in them.

```
[root@master-node1 yaml]# kubectl describe pod mysql-7bc745d4f8-wnr2p -n techtracker
Name:           mysql-7bc745d4f8-wnr2p
Namespace:      techtracker
Priority:      0
Node:          worker-node1/192.168.0.137
Start Time:    Sat, 24 Dec 2022 07:12:18 -0500
Labels:         app=mysql
               pod-template-hash=7bc745d4f8
Annotations:   <none>
Status:        Running
IP:            10.32.0.2
IPs:
  IP:          10.32.0.2
Controlled By: ReplicaSet/mysql-7bc745d4f8
Containers:
  mysql:
    Container ID:  docker://22403cf8a84b58f704e4049ca8b318b25babf5f667b4fcbbaa2958c8
    Image:         mysql:5.6
    Image ID:     docker-pullable://mysql@sha256:20575ecebe6216036d25dab5903808211f
    Port:          <none>
    Host Port:    <none>
    State:        Running
      Started:   Sat, 24 Dec 2022 07:12:54 -0500
    Ready:        True
    Restart Count: 0
    Environment:
      MYSQL_ROOT_PASSWORD: password
    Mounts:
```

Create Secret

```
kubectl create secret generic mysqlsecret --from-literal=mysqlpassword=redhat  
-n techtracker  
kubectl get secret mysqlsecret -n techtracker  
kubectl describe secret mysqlsecret -n techtracker
```

```
[root@master-node1 yaml]# kubectl create secret generic mysqlsecret --from-literal=mysqlpassword=redhat -n techtracker  
secret/mysqlsecret created  
[root@master-node1 yaml]# kubectl get secret mysqlsecret -n techtracker  
NAME      TYPE   DATA  AGE  
mysqlsecret  Opaque  1    41s  
[root@master-node1 yaml]# kubectl describe secret mysqlsecret -n techtracker  
Name:         mysqlsecret  
Namespace:    techtracker  
Labels:       <none>  
Annotations:  <none>  
  
Type:  Opaque  
  
Data  
====  
mysqlpassword: 6 bytes
```

6 bytes means 6 digits (redhat)

Create deployment using secret values

```
apiVersion: apps/v1  
kind: Deployment  
metadata:  
  name: mysql  
  namespace: techtracker  
spec:  
  selector:  
    matchLabels:  
      app: mysql  
  template:  
    metadata:  
      labels:  
        app: mysql  
    spec:  
      containers:  
        - image: mysql:5.6  
          name: mysql  
          env:  
            - name: MYSQL_ROOT_PASSWORD  
              valueFrom:  
                secretKeyRef:  
                  name: mysqlsecret  
                  key: mysqlpassword
```

```
kubectl apply -f resource_valuefromsecret.yaml  
kubectl get all -n techtracker
```

```
[root@master-node1 yaml]# kubectl apply -f resource_valuefromsecret.yaml  
deployment.apps/mysql created  
[root@master-node1 yaml]# kubectl get all -n techtracker  
NAME                                     READY   STATUS    RESTARTS   AGE  
pod/mysql-74874b5587-tnv47   0/1     Pending   0          9s  
  
NAME                           READY   UP-TO-DATE   AVAILABLE   AGE  
deployment.apps/mysql   0/1     1           0          9s  
  
NAME                         DESIRED  CURRENT  READY   AGE  
replicaset.apps/mysql-74874b5587  1        1        0      9s
```

```
kubectl describe pod mysql-74874b5587-tnv47 -n techtracker
```

```
[root@master-node1 ~]# kubectl describe pod mysql-74874b5587-tnv47 -n techtracker  
Name:           mysql-74874b5587-tnv47  
Namespace:      techtracker  
Priority:      0  
Node:          <none>  
Labels:         app=mysql  
               pod-template-hash=74874b5587  
Annotations:    <none>  
Status:         Pending  
IP:  
IPs:          <none>  
Controlled By: ReplicaSet/mysql-74874b5587  
Containers:  
  mysql:  
    Image:      mysql:5.6  
    Port:       <none>  
    Host Port: <none>  
    Environment:  
      MYSQL_ROOT_PASSWORD: <set to the key 'mysqlpassword' in secret 'mysqlsecret'> Optional: false  
    Mounts:  
      /var/run/secrets/kubernetes.io/serviceaccount from default-token-8thdx (ro)  
Conditions:  
  Type      Status  
  PodScheduled  False
```

env variable value are not displayed now when we describe pod

Person having Rights to edit the secret can only view the value stored in the key .

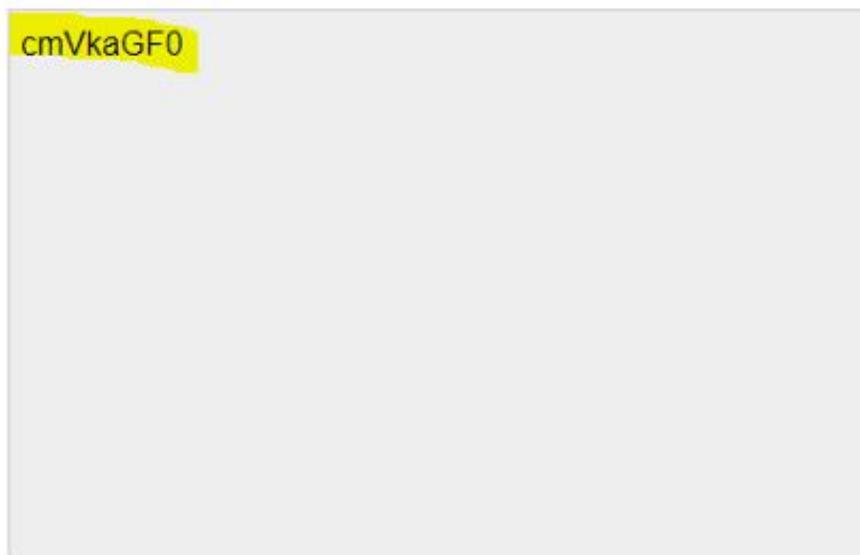
```
kubectl edit secret mysqlsecret -n techtracker
```

```
# Please edit the object below. Lines beginning with '#'
# and an empty file will abort the edit. If an
# reopened with the relevant failures.
#
apiVersion: v1
data:
  mysqlpassword: cmVkaGF0
kind: Secret
metadata:
  creationTimestamp: "2022-12-24T12:48:38Z"
  managedFields:
    - apiVersion: v1
      fieldsType: FieldsV1
```

By default Base64 encoded value stored in secret (Highlighted in Yellow above)

Decode from Base64 format

Simply enter your data then push the decode button.



The screenshot shows a web-based Base64 decoder. At the top, there is a text input field containing the string "cmVkaGF0". Below the input field is a large, empty gray area intended for the decoded output. The interface is clean and minimalist.

ⓘ For encoded binaries (like images, documents, etc.) use the [encoder](#).

UTF-8 ▼ Source character set.

Decode each line separately (useful for when you have multi-line text).

Live mode OFF Decodes in real-time as you type or paste.

< DECODE > Decodes your data into the area below.

redhat

we can create secret using yaml also but value we have to put as Base64 encoded value.

```
apiVersion: apps/v1
kind: Secret
metadata:
  name: sample-secret
  namespace: techtracker
type: opaque
data:
  mysqlrootpassword: a3vhFjsakD3jKJ34HJ█
```

Create Configmaps

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: mysql
  namespace: techtracker
spec:
  selector:
    matchLabels:
      app: mysql
  template:
    metadata:
      labels:
        app: mysql
    spec:
      containers:
        - image: mysql:5.6
          name: mysql
          env:
            - name: MYSQL_ROOT_PASSWORD
              valueFrom:
                configMapKeyRef:
                  name: mysqlcm
                  key: mysqlpassword
  kubectl create configmap mysqlcm --from-literal=mysqlpassword=redhat -n techtracker
  kubectl apply -f resource_valuefromConfigmap.yaml
  kubectl get all -n techtracker -o wide
  kubectl describe pod mysql-859699c87c-9hkp7 -n techtracker
```

```
[root@master-node1 yaml]# kubectl apply -f resource_valuefromConfigmap.yaml
deployment.apps/mysql created
[root@master-node1 yaml]# kubectl get all -n techtracker -o wide
NAME                                READY   STATUS    RESTARTS   AGE     IP          NODE      NOMINATED NODE   READINESS
pod/mysql-859699c87c-9hkp7   0/1    Pending   0          4s     <none>    <none>    <none>    <none>
NAME                                READY   UP-TO-DATE   AVAILABLE   AGE     CONTAINERS   IMAGES   SELECTOR
deployment.apps/mysql   0/1    1           0          4s     mysql        mysql:5.6   app=mysql
NAME                                DESIRED  CURRENT  READY     AGE     CONTAINERS   IMAGES   SELECTOR
replicaset.apps/mysql-859699c87c  1        1        0         4s     mysql        mysql:5.6   app=mysql,pod-
[root@master-node1 yaml]# kubectl describe pod mysql-859699c87c-9hkp7 -n techtracker
Name:           mysql-859699c87c-9hkp7
Namespace:      techtracker
Priority:      0
Node:          <none>
Labels:         app=mysql
                pod-template-hash=859699c87c
Annotations:    <none>
Status:        Pending
IP:
IPs:          <none>
Controlled By: ReplicaSet/mysql-859699c87c
Containers:
  mysql:
    Image:      mysql:5.6
    Port:       <none>
    Host Port: <none>
    Environment:
      MYSQL_ROOT_PASSWORD: <set to the key 'mysqlpassword' of config map 'mysqlcm'> Optional: false
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from default-token-8thdx (ro)

```

```
kubectl describe configmap mysqlcm -n techtracker
```

In configmap we can see value using describer

but in Secret we need edit rights to see value

```
[root@master-node1 yaml]# kubectl describe configmap mysqlcm -n techtracker
Name:           mysqlcm
Namespace:      techtracker
Labels:         <none>
Annotations:    <none>

Data
====
mysqlpassword: redhat
```

13. Resource Quota

1. How many Pods,Service,Deployment,Secret,ConfigMap we can create in a particular Namespace , Its defined under:

Namespace Quota – Project quota – object Quota – resource quota.yaml

2. How much Resource the object can consume — project quota
3. How much resource a Pod can use and How much resource a container can use inside a Pod .

Create Resource Quota to Limit Pod Count as 4

```
apiVersion: v1
kind: ResourceQuota
metadata:
  name: pods-high
  namespace: techtracker
spec:
  hard:
    pods: "4"
kubectl create -f resource_quota.yaml
kubectl get resourcequota -n techtracker
```

```
[root@master-node1 yaml]# kubectl create -f resource_quota.yaml
resourcequota/pods-high created
[root@master-node1 yaml]# kubectl get resourcequota -n techtracker
NAME      AGE     REQUEST      LIMIT
pods-high  26s    pods: 0/4
```

Now if we try to create 10 Pods using deployment , Only 4 will be created

```
kubectl apply -f deployment.yaml  
kubectl get all -n techtracker
```

```
[root@master-node1 yaml]# kubectl get all -n techtracker  
NAME                                         READY   STATUS    RESTARTS   AGE  
pod/techtracker-deployment-6b474476c4-jcjxr  1/1     Running   0          93s  
pod/techtracker-deployment-6b474476c4-v2qc8  1/1     Running   0          93s  
pod/techtracker-deployment-6b474476c4-vmttk  1/1     Running   0          93s  
pod/techtracker-deployment-6b474476c4-w6nbp  1/1     Running   0          93s  
  
NAME                           READY   UP-TO-DATE   AVAILABLE   AGE  
deployment.apps/techtracker-deployment  4/10   4           4          93s  
  
NAME                               DESIRED   CURRENT   READY   AGE  
replicaset.apps/techtracker-deployment-6b474476c4  10       4         4        93s
```

14. Request and Limit

Requests:memory:900Mi means 900mb minimum is required at Node to launch this Pod

Limits:memory:10Gi means 10gb max can be used by container while running

```
resources:  
  requests:  
    memory: "900Mi"  
    cpu: "500m"  
  limits:  
    memory: "10Gi"  
    cpu: "500m"
```

15. Edit any kubernetes resource with option -o yaml & Delete the pods after editing replicaset

```
kubectl edit replicaset.apps/new-replica-set -o yaml  
kubectl delete pod new-replica-set-7c2r9  
kubectl delete pod/new-replica-set-bcwgg  
kubectl delete pod/new-replica-set-nxc9r  
kubectl delete pod/new-replica-set-sj77c
```

16. Scale replicaSet

```
kubectl scale replicaset.apps/new-replica-set --replicas=5  
kubectl scale replicaset.apps/new-replica-set --replicas=2
```

```
controlplane ~ ➔ kubectl scale replicaset.apps/new-replica-set --replicas=5
replicaset.apps/new-replica-set scaled

controlplane ~ ➔ kubectl get all
NAME                      READY   STATUS    RESTARTS   AGE
pod/new-replica-set-q7n56  1/1     Running   0          13m
pod/new-replica-set-xzvc5  1/1     Running   0          13m
pod/new-replica-set-vjdpd  1/1     Running   0          12m
pod/new-replica-set-9wmmp  1/1     Running   0          12m
pod/new-replica-set-rch75  1/1     Running   0          10s

NAME            TYPE        CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
service/kubernetes   ClusterIP   10.43.0.1      <none>           443/TCP      51m

NAME              DESIRED   CURRENT   READY   AGE
replicaset.apps/new-replica-set  5         5         5       44m

controlplane ~ ➔ kubectl scale replicaset.apps/new-replica-set --replicas=2
replicaset.apps/new-replica-set scaled

controlplane ~ ➔ kubectl get all
NAME                      READY   STATUS    RESTARTS   AGE
pod/new-replica-set-xzvc5  1/1     Running   0          14m
pod/new-replica-set-vjdpd  1/1     Running   0          14m
pod/new-replica-set-9wmmp  1/1     Terminating   0          14m
pod/new-replica-set-q7n56  1/1     Terminating   0          14m
pod/new-replica-set-rch75  1/1     Terminating   0          104s

NAME            TYPE        CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
service/kubernetes   ClusterIP   10.43.0.1      <none>           443/TCP      53m

NAME              DESIRED   CURRENT   READY   AGE
replicaset.apps/new-replica-set  2         2         2       45m
```

17. Creating static pod without replicaset and deployment

pod.yaml

```
apiVersion: v1
kind: Pod
metadata:
  name: redis
  namespace: finance
  labels:
    app: redis
spec:
  containers:
    - name: redis
      image: redis
~  
kubectl apply -f pod.yaml  
kubectl get pods -n finance
```

```
controlplane ~ ➔ kubectl get pods -n finance
NAME      READY   STATUS    RESTARTS   AGE
payroll   1/1     Running   0          12m
redis     1/1     Running   0          2m2s
```

18. Find a Pod in all namespaces

```
kubectl get pods --all-namespaces
```

```
controlplane ~ ➔ kubectl get pods --all-namespaces
NAMESPACE      NAME           READY   STATUS    RESTARTS   AGE
kube-system   local-path-provisioner-7b7dc8d6f5-8q988  1/1     Running   0          22m
kube-system   coredns-b96499967-z24x5      1/1     Running   0          22m
kube-system   helm-install-traefik-crd-m988z  0/1     Completed  0          22m
kube-system   helm-install-traefik-9qtgt  0/1     Completed  1          22m
kube-system   traefik-7cd4fcff68-xzq6w    1/1     Running   0          19m
kube-system   svclb-traefik-mkmkm     2/2     Running   0          19m
marketing     redis-db        1/1     Running   0          16m
dev           redis-db        1/1     Running   0          16m
marketing     blue            1/1     Running   0          16m
finance       payroll         1/1     Running   0          16m
manufacturing red-app        1/1     Running   0          16m
research      dna-1          0/1     CrashLoopBackOff  6 (4m59s ago)  16m
kube-system   metrics-server-668d979685-xzgzd  1/1     Running   0          22m
research      dna-2          0/1     CrashLoopBackOff  6 (4m51s ago)  16m
finance       redis          1/1     Running   0          5m25s
```

```
kubectl get pods --all-namespaces -l name=blue
```

```
controlplane ~ ➔ kubectl get pods --all-namespaces -l name=blue
NAMESPACE      NAME      READY      STATUS      RESTARTS      AGE
marketing      blue      1/1       Running     0            20m
```

```
kubectl get pods nginx-pod -o yaml
```

```
controlplane ~ ➔ kubectl get pods nginx-pod -o yaml
apiVersion: v1
kind: Pod
metadata:
  creationTimestamp: "2022-12-26T10:58:56Z"
  labels:
    run: nginx-pod
    name: nginx-pod
    namespace: default
    resourceVersion: "728"
    uid: e485c357-c5b3-4410-a0ac-4d94456bd69c
spec:
  containers:
  - image: nginx:alpine
    imagePullPolicy: IfNotPresent
    name: nginx-pod
  
```

19. Creating pod using Imperative command using dr-run=client

```
kubectl run nginx-pod --image=nginx:alpine -o yaml --dry-run=client
```

```
controlplane ~ ➔ kubectl run nginx-pod --image=nginx:alpine -o yaml --dry-run=client
apiVersion: v1
kind: Pod
metadata:
  creationTimestamp: null
  labels:
    run: nginx-pod
    name: nginx-pod
spec:
  containers:
  - image: nginx:alpine
    name: nginx-pod
    resources: {}
  dnsPolicy: ClusterFirst
  restartPolicy: Always
status: {}
```

```
kubectl run redis --image=redis:alpine -l tier=db -o yaml --dry-run=client
```

```
controlplane ~ ➔ kubectl run redis --image=redis:alpine -l tier=db -o yaml --dry-run=client
apiVersion: v1
kind: Pod
metadata:
  creationTimestamp: null
  labels:
    tier: db
    name: redis
spec:
  containers:
  - image: redis:alpine
    name: redis
    resources: {}
  dnsPolicy: ClusterFirst
  restartPolicy: Always
status: {}
```

20. Creating a ClusterIP service using kubectl expose

```
kubectl expose pod redis --port=6379 --name=redis-service --dry-run=client  
kubectl get svc  
kubectl expose pod redis --port=6379 --name=redis-service  
kubectl get svc
```

```
controlplane ~ ➔ kubectl expose pod redis --port=6379 --name=redis-service --dry-run=client  
service/redis-service exposed (dry run)  
  
controlplane ~ ➔ kubectl get svc  
NAME      TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE  
kubernetes   ClusterIP   10.43.0.1    <none>        443/TCP     20m  
  
controlplane ~ ➔ kubectl expose pod redis --port=6379 --name=redis-service  
service/redis-service exposed  
  
controlplane ~ ➔ kubectl get svc  
NAME      TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE  
kubernetes   ClusterIP   10.43.0.1    <none>        443/TCP     21m  
redis-service   ClusterIP   10.43.53.213  <none>        6379/TCP    5s
```

21. Creating deployment using Imperative command

```
kubectl create deployment webapp --image=kodekloud/webapp-color --replicas=3  
kubectl get all
```

```
controlplane ~ ➔ kubectl create deployment webapp --image=kodekloud/webapp-color --replicas=3
deployment.apps/webapp created

controlplane ~ ➔ kubectl get all
NAME                      READY   STATUS    RESTARTS   AGE
pod/nginx-pod              1/1     Running   0          21m
pod/redis                  1/1     Running   0          13m
pod/webapp-79c4cc587b-swzml 1/1     Running   0          14s
pod/webapp-79c4cc587b-ldn7q 1/1     Running   0          14s
pod/webapp-79c4cc587b-bbhpd 1/1     Running   0          14s

NAME            TYPE        CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
service/kubernetes  ClusterIP   10.43.0.1      <none>           443/TCP     29m
service/redis-service ClusterIP  10.43.53.213  <none>           6379/TCP    8m15s

NAME          READY   UP-TO-DATE   AVAILABLE   AGE
deployment.apps/webapp  3/3     3           3           14s

NAME          DESIRED   CURRENT   READY   AGE
replicaset.apps/webapp-79c4cc587b  3         3         3       14s
```

22. Create a pod with container port 8080

```
kubectl run custom-nginx --image=nginx --port=8080
```

```
kubectl describe pod custom-nginx | grep -i port
```

```
controlplane ~ ➔ kubectl get pods -o wide
NAME                      READY   STATUS    RESTARTS   AGE      IP           NODE
nginx-pod                 1/1     Running   0          25m     10.42.0.9   controlplane
redis                     1/1     Running   0          16m     10.42.0.10  controlplane
webapp-79c4cc587b-swzml  1/1     Running   0          3m51s   10.42.0.13  controlplane
webapp-79c4cc587b-ldn7q  1/1     Running   0          3m51s   10.42.0.12  controlplane
webapp-79c4cc587b-bbhpd  1/1     Running   0          3m51s   10.42.0.11  controlplane
custom-nginx               1/1     Running   0          24s     10.42.0.14  controlplane

controlplane ~ ➔ kubectl describe pod custom-nginx | grep -i port
  Port:                8080/TCP
  Host Port:           0/TCP
```

23. Create ns, deployment using Imperative command

```
kubectl create ns dev-ns  
kubectl create deployment redis-deploy --image=redis --replicas=2 -n dev-ns  
kubectl get all -n dev-ns
```

```
controlplane ~ ➔ kubectl get all -n dev-ns  
NAME                      READY   STATUS    RESTARTS   AGE  
pod/redis-deploy-f9b7599c5-vjqgq2   1/1     Running   0          26s  
pod/redis-deploy-f9b7599c5-c97g9   1/1     Running   0          26s  
  
NAME                      READY   UP-TO-DATE   AVAILABLE   AGE  
deployment.apps/redis-deploy   2/2     2           2          27s  
  
NAME                      DESIRED  CURRENT  READY   AGE  
replicaset.apps/redis-deploy-f9b7599c5   2        2        2       27s
```

24. Create a httpd pod and expose it on port 80

```
kubectl run httpd --image=httpd:alpine --port=80 --expose  
kubectl get all
```

```
controlplane ~ ➔ kubectl get all
NAME                      READY   STATUS    RESTARTS   AGE
pod/nginx-pod              1/1     Running   0          47m
pod/redis                  1/1     Running   0          38m
pod/webapp-79c4cc587b-swzml 1/1     Running   0          25m
pod/webapp-79c4cc587b-ldn7q 1/1     Running   0          25m
pod/webapp-79c4cc587b-bbhpd 1/1     Running   0          25m
pod/custom-nginx            1/1     Running   0          22m
pod/httpd                  1/1     Running   0          31s

NAME           TYPE        CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
service/kubernetes ClusterIP  10.43.0.1      <none>        443/TCP     54m
service/redis-service ClusterIP  10.43.53.213  <none>        6379/TCP    33m
service/httpd       ClusterIP  10.43.211.104 <none>        80/TCP      31s

NAME          READY   UP-TO-DATE   AVAILABLE   AGE
deployment.apps/webapp  3/3      3           3           25m

NAME          DESIRED   CURRENT   READY   AGE
replicaset.apps/webapp-79c4cc587b  3         3         3         25m
```

25. Enter Sleep with timer in container command

```
apiVersion: v1
kind: Pod
metadata:
  name: ubuntu-sleeper-2
spec:
  containers:
  - name: ubuntu
    image: ubuntu
    command:
    - sleep
    - "5000"
```

```
controlplane ~ ➔ vi ubuntu-sleeper-2.yaml

controlplane ~ ➔ kubectl apply -f ubuntu-sleeper-2.yaml
pod/ubuntu-sleeper-2 created

controlplane ~ ➔ kubectl get pods
NAME           READY   STATUS    RESTARTS   AGE
ubuntu-sleeper   1/1     Running   0          5m29s
ubuntu-sleeper-2  1/1     Running   0          10s
```

26. Create ConfigMap using Command

```
create configmap webapp-config-map --from-literal=APP_COLOR=darkblue
get cm
describe cm webapp-config-map
```

```
controlplane ~ ➔ kubectl create configmap webapp-config-map --from-literal=APP_COLOR=darkblue
configmap/webapp-config-map created

controlplane ~ ➔ kubectl get cm
NAME        DATA   AGE
kube-root-ca.crt  1     16m
db-config      3     3m1s
webapp-config-map 1     9s

controlplane ~ ➔ kubectl describe cm webapp-config-map
Name:         webapp-config-map
Namespace:    default
Labels:       <none>
Annotations: <none>

Data
====
APP_COLOR:
-----
darkblue

BinaryData
=====
```

```
spec:  
  containers:  
    - env:  
      - name: APP_COLOR  
        valueFrom:  
          configMapKeyRef:  
            name: webapp-config-map  
            key: APP_COLOR
```

27. Create Secret having multiple key using command

```
kubectl create secret generic db-secret --from-literal=DB_Host=sql01 --from-literal=DB_User=root --from-literal=DB_Password=password123  
kubectl describe secret db-secret
```

```
controlplane ~ ➔ kubectl describe secret db-secret  
Name:         db-secret  
Namespace:    default  
Labels:       <none>  
Annotations:  <none>  
  
Type:  Opaque  
  
Data  
====  
DB_Host:      5 bytes  
DB_Password:  11 bytes  
DB_User:      4 bytes
```

```
spec:  
  containers:  
    - env:  
        - name: DB_Host  
          valueFrom:  
            secretKeyRef:  
              name: db-secret  
              key: DB_Host  
        - name: DB_User  
          valueFrom:  
            secretKeyRef:  
              name: db-secret  
              key: DB_User  
        - name: DB_Password  
          valueFrom:  
            secretKeyRef:  
              name: db-secret  
              key: DB_Password  
    - image: kodekloud/simple-webapp-mysql  
      imagePullPolicy: Always  
      name: webapp  
  kubectl apply -f /tmp/kubectl-edit-1663824269.yaml
```

28. Find out user used to execute the sleep process within the `ubuntu-sleeper` pod ?

```
kubectl exec ubuntu-sleeper -it -- whoami
```

```
controlplane ~ ✘ kubectl exec ubuntu-sleeper -it -- whoami  
root
```

29. Change default user to run command in container using securityContext

```
apiVersion: v1
kind: Pod
metadata:
  creationTimestamp: "2022-12-27T07:48:34Z"
  name: ubuntu-sleeper
  namespace: default
  resourceVersion: "718"
  uid: 9f52cb0e-475b-4174-a657-a3a4994822e9
spec:
  securityContext:
    runAsUser: 1010
  containers:
    - command:
      - sleep
      - "4800"
    image: ubuntu
```

30. save running Pod output to yaml file

```
kubectl get pod ubuntu-sleeper -o yaml > /tmp/test1.yaml
```

31. securityContext capability

```
spec:  
  containers:  
    - command:  
      - sleep  
      - "4800"  
    image: ubuntu  
    imagePullPolicy: Always  
    name: ubuntu  
    securityContext:  
      capabilities:  
        add: ["SYS_TIME"]  
    resources: {}
```

32. Multiple capabilities in Array

```
containers:  
  - command:  
    - sleep  
    - "4800"  
  image: ubuntu  
  imagePullPolicy: Always  
  name: ubuntu  
  securityContext:  
    capabilities:  
      add: ["SYS_TIME", "NET_ADMIN"]  
  resources: {}
```

33. Check existing serviceaccounts

```
kubectl get serviceaccounts  
kubectl describe sa default  
kubectl create serviceaccount dashboard-sa
```

```
controlplane ~ ➔ kubectl get serviceaccounts  
NAME      SECRETS   AGE  
default    0          4m36s  
dev       0          51s
```

```
controlplane ~ ➔ kubectl describe sa default
Name:                  default
Namespace:             default
Labels:                <none>
Annotations:           <none>
Image pull secrets:   <none>
Mountable secrets:    <none>
Tokens:                <none>
Events:                <none>
```

34. Create token using serviceAccount

```
kubectl create token dashboard-sa
```

```
controlplane /var/rbac ➔ kubectl create token dashboard-sa
eyJhbGciOiJSUzI1NiIsImtpZCI6ImMwdjRIBkRiaFVEOTlBS244VHR0YnRta3RuNzJqdXkxMkJ6MW1dFZjVkJoxNjcyMTM3M
WQiOlsiaHR0cHM6Ly9rdWJlc5ldGVzLmRlZmF1bHQuc3ZjLmNsdxN0ZXIubG9jYwWiLCJrM3MiXSwiZKhwIjoxNjcyMTM3M
zE2LCJpYXQiOjE2NzIxMzM3MTYsImIzcyI6Imh0dBzOi8va3VizXJuZXR1cy5kZWZhdx0LnN2Yy5jbHVzdGVyLmxvY2FsI
iwiia3VizXJuZXR1cy5pbbyI6eyJuYW1lc3BhY2UiOijkZWZhdx0Iiwiic2VydmljZWFjY291bnQiOnsibmFtZSI6ImRhc2hib
2FyZC1zYSIsInVpZCI6IjBlZDUyMzk4LTRizGUtNDQyOC1iMmUyLTthhYzQxZWYxOWEzYyJ9fSwibmJmIjoxNjcyMTMzNzE2L
CJzdWIiOijzeXN0ZW6c2VydmljZWFjY291bnQ6ZGVmYXVsDpkYXNoYm9hcmQtc2EifQ.YAk--y8jo-BwwNa9jxP5spx2Y4
_TSQlUV0-mhOh24ck678ZigZHTUGwnuZCn9S4bFCEmOhoJIt490YH4YxJC6Gnjj--DFXnejmFsPZKqum9WRhpHXNyNy3nFA
G2piHO1dWB3jokHsjbRNPL84gzX81LbDVIIHyZmkN7pz4ZZ_7f7NKLWyG4kRkFOArFU1dp4KM_hvdBOGu1tvQRD_vNB0-DMr
LrwugQ2cgdJVK90YzC11MxRBc1Q-kM1_R82ZpYHwGTrGyPxh4hizNIUxURXeDG1iZ8fMHp4aLBkB81ecNtbd7XBeukOnH61Y
ny6sxNclh2GXD_rqaldcou0BybSA
```

```
metadata:
  creationTimestamp: null
  labels:
    name: web-dashboard
spec:
  serviceAccountName: dashboard-sa
  containers:
  - env:
```

35. Check logs of a pod

```
kubectl -n elastic-stack logs kibana
```

```
controlplane ~ ➔ kubectl -n elastic-stack logs kibana
{"type": "log", "@timestamp": "2022-12-27T12:32:40Z", "tags": ["status", "plugin:kib", "pid": 1, "state": "green", "message": "Status changed from uninitialized to green"}, "error": false, "file": "kibana", "line": 1, "offset": 1, "source": "elasticsearch", "type": "log", "version": 1}
{"type": "log", "@timestamp": "2022-12-27T12:32:40Z", "tags": ["status", "plugin:kib", "pid": 1, "state": "green", "message": "Status changed from uninitialized to green"}, "error": false, "file": "kibana", "line": 1, "offset": 1, "source": "elasticsearch", "type": "log", "version": 1}
```

36. Liveness, Readiness, and Startup Probes

Startup Probe

A startup probe verifies whether the application within a container is started. Startup probes run before any other probe, and, unless it finishes successfully, disables other probes. If a container fails its startup probe, then the container is killed and follows the pod's `restartPolicy`.

This type of probe is only executed at startup, unlike readiness probes, which are run periodically.

The startup probe is configured in the `spec.containers.startupProbe` attribute of the pod configuration.

Readiness Probe

Readiness probes determine whether or not a container is ready to serve requests. If the readiness probe returns a failed state, then

Kubernetes removes the IP address for the container from the endpoints of all Services.

Developers use readiness probes to instruct Kubernetes that a running container should not receive any traffic. This is useful when waiting for an application to perform time-consuming initial tasks, such as establishing network connections, loading files, and warming caches.

The readiness probe is configured in
the `spec.containers.readinessprobe` attribute of the pod configuration.

Liveness Probe

Liveness probes determine whether or not an application running in a container is in a `healthy` state. If the liveness probe detects an unhealthy state, then Kubernetes kills the container and tries to redeploy it.

The liveness probe is configured in
the `spec.containers.livenessprobe` attribute of the pod configuration.

```
apiVersion: v1
kind: Pod
metadata:
  creationTimestamp: "2021-08-01T04:55:35Z"
  labels:
    name: simple-webapp
    name: simple-webapp-2
    namespace: default
spec:
  containers:
    - env:
      - name: APP_START_DELAY
```

```
    value: "80"
  image: kodekloud/webapp-delayed-start
  imagePullPolicy: Always
  name: simple-webapp
  ports:
  - containerPort: 8080
    protocol: TCP
  readinessProbe:
    httpGet:
      path: /ready
      port: 8080
~
```

To recreate the pod, run the command:

```
kubectl replace -f simple-webapp-2.yaml --force
controlplane ~ ➔ kubectl replace -f new.yaml --force
pod "simple-webapp-2" deleted
pod/simple-webapp-2 replaced
```

```
apiVersion: v1
kind: Pod
metadata:
  labels:
    name: simple-webapp
    name: simple-webapp-1
    namespace: default
spec:
  containers:
  - env:
    - name: APP_START_DELAY
      value: "80"
    image: kodekloud/webapp-delayed-start
    imagePullPolicy: Always
    name: simple-webapp
    ports:
    - containerPort: 8080
      protocol: TCP
    readinessProbe:
      httpGet:
        path: /ready
        port: 8080
    livenessProbe:
      httpGet:
        path: /live
        port: 8080
```

```
periodSeconds: 1
initialDelaySeconds: 80
```

37. check memory of nodes in cluster

```
kubectl top node
kubectl top pods
```

```
controlplane kubernetes-metrics-server on master ➔ kubectl top node
NAME          CPU(cores)   CPU%    MEMORY(bytes)   MEMORY%
controlplane   317m        0%      1216Mi         0%
node01         36m        0%      292Mi          0%
```

```
controlplane kubernetes-metrics-server on master ➔ kubectl top pods
NAME          CPU(cores)   MEMORY(bytes)
elephant      0m          33Mi
lion          0m          19Mi
rabbit        0m          252Mi
```

38. *InitContainer* yaml

```
spec:
  containers:
  - command:
    - sh
    - -c
    - echo The app is running! && sleep 3600
  image: busybox:1.28
  imagePullPolicy: IfNotPresent
  name: green-container-1
  resources: {}
  terminationMessagePath: /dev/termination-log
  terminationMessagePolicy: File
  volumeMounts:
  - mountPath: /var/run/secrets/kubernetes.io/serviceaccount
    name: kube-api-access-2dz76
    readOnly: true
  dnsPolicy: ClusterFirst
  enableServiceLinks: true
  initContainers:
  - command:
```

```
- sh
- -c
- sleep 5
image: busybox
imagePullPolicy: Always
name: init-my-service
resources: {}
terminationMessagePath: /dev/termination-log
terminationMessagePolicy: File
volumeMounts:
- mountPath: /var/run/secrets/kubernetes.io/serviceaccount
  name: kube-api-access-2dz76
  readOnly: true
nodeName: controlplane
```

39. Find Resources with specific Multiple labels

```
kubectl get pods -l env=dev
kubectl get pods -l bu=finance
kubectl get all -l env=prod
kubectl get pods -l env=prod,bu=finance,tier=frontend
```

controlplane ~ ➔ kubectl get pods -l env=dev				
NAME	READY	STATUS	RESTARTS	AGE
app-1-wtcvd	1/1	Running	0	2m4s
db-1-h15mp	1/1	Running	0	2m4s
app-1-s8nnq	1/1	Running	0	2m4s
db-1-dhdx7	1/1	Running	0	2m4s
db-1-rzw77	1/1	Running	0	2m4s
app-1-r4xgq	1/1	Running	0	2m4s
db-1-57hrt	1/1	Running	0	2m4s

```
controlplane ~ ➔ kubectl get pods -l bu=finance
NAME        READY   STATUS    RESTARTS   AGE
app-1-wtcvd 1/1     Running   0          7m48s
app-1-s8nnq  1/1     Running   0          7m48s
app-1-zzxdf  1/1     Running   0          7m47s
db-2-srp77   1/1     Running   0          7m48s
app-1-r4xgq  1/1     Running   0          7m48s
auth         1/1     Running   0          7m48s

controlplane ~ ➔ kubectl get all -l env=prod
NAME        READY   STATUS    RESTARTS   AGE
pod/app-1-zzxdf 1/1     Running   0          19m
pod/db-2-srp77  1/1     Running   0          19m
pod/app-2-w2zph 1/1     Running   0          19m
pod/auth       1/1     Running   0          19m

NAME           TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
service/app-1  ClusterIP  10.43.87.35   <none>        3306/TCP   19m

NAME           DESIRED  CURRENT  READY   AGE
replicaset.apps/db-2  1        1        1      19m
replicaset.apps/app-2 1        1        1      19m

controlplane ~ ➔ kubectl get pods -l env=prod,bu=finance,tier=frontend
NAME        READY   STATUS    RESTARTS   AGE
app-1-zzxdf 1/1     Running   0          24m
```

40. Create Job using yaml and Imperative command

```
apiVersion: batch/v1
kind: Job
metadata:
  name: throw-dice-job
spec:
  backoffLimit: 15 # This is so the job does not quit before it succeeds.
  template:
    spec:
      containers:
        - name: throw-dice
          image: kodekloud/throw-dice
      restartPolicy: Never
  kubectl create job throw-dice-job --image=kodekloud/throw-dice --dry-run=client -o yaml > throw-dice-job.yaml
```

```
kubectl create job throw-dice-job --image=kodekloud/throw-dice -o yaml >
throw-dice-job.yaml
apiVersion: batch/v1
kind: Job
metadata:
  name: throw-dice-job
spec:
  completions: 3
  backoffLimit: 25 # This is so the job does not quit before it succeeds.
  template:
    spec:
      containers:
        - name: throw-dice
          image: kodekloud/throw-dice
      restartPolicy: Never
```

41. Describe job to check success and failed logs in events

```
kubectl describe job throw-dice-job
```

```
Start Time:      Wed, 28 Dec 2022 08:29:05 +0000
Completed At:    Wed, 28 Dec 2022 08:29:39 +0000
Duration:       34s
Pods Statuses:   0 Active (0 Ready) / 3 Succeeded / 2 Failed
Pod Template:
  Labels: controller-uid=1fdf6228-2757-4087-be9f-02340ee96a59
           job-name=throw-dice-job
  Containers:
    throw-dice:
      Image:      kodekloud/throw-dice
      Port:       <none>
      Host Port:  <none>
      Environment: <none>
      Mounts:     <none>
      Volumes:    <none>
  Events:
    Type  Reason     Age   From            Message
    ----  ----     --   --              --
    Normal SuccessfulCreate 6m12s  job-controller  Created pod: throw-dice-job-l29v8
    Normal SuccessfulCreate 6m8s   job-controller  Created pod: throw-dice-job-lvjjc
    Normal SuccessfulCreate 5m53s  job-controller  Created pod: throw-dice-job-6kzqd
    Normal SuccessfulCreate 5m48s  job-controller  Created pod: throw-dice-job-blx8l
    Normal SuccessfulCreate 5m43s  job-controller  Created pod: throw-dice-job-2fsxp
    Normal Completed     5m38s  job-controller  Job completed
```

42. Job finishes in 1 sec using 3 parallel job parallelism: 3

```
apiVersion: batch/v1
kind: Job
metadata:
  name: throw-dice-job
spec:
  completions: 3
  parallelism: 3
  backoffLimit: 25 # This is so the job does not quit before it succeeds.
  template:
    spec:
      containers:
        - name: throw-dice
          image: kodekloud/throw-dice
          restartPolicy: Never
~  
kubectl describe job throw-dice-job
```

```
Completion Mode: NonIndexed
Start Time:       Wed, 28 Dec 2022 08:40:52 +0000
Completed At:     Wed, 28 Dec 2022 08:41:06 +0000
Duration:         14s
Pods Statuses:   0 Active (0 Ready) / 3 Succeeded / 2 Failed
Pod Template:
  Labels: controller-uid=f81f3ad7-c7c4-4dd5-928f-759bf963f2bb
            job-name=throw-dice-job
  Containers:
    throw-dice:
      Image:      kodekloud/throw-dice
      Port:       <none>
      Host Port:  <none>
      Environment: <none>
      Mounts:     <none>
      Volumes:    <none>
  Events:
    Type      Reason          Age    From           Message
    ----      ----          ----   ----          -----
    Normal   SuccessfulCreate  18s   job-controller  Created pod: throw-dice-job-5vz2r
    Normal   SuccessfulCreate  18s   job-controller  Created pod: throw-dice-job-cbh5m
    Normal   SuccessfulCreate  18s   job-controller  Created pod: throw-dice-job-jm8mw
    Normal   SuccessfulCreate  12s   job-controller  Created pod: throw-dice-job-ks2gs
    Normal   SuccessfulCreate  10s   job-controller  Created pod: throw-dice-job-qr4lq
    Normal   Completed        4s    job-controller  Job completed
```

43. Schedule cron job in yaml

```
apiVersion: batch/v1
kind: CronJob
metadata:
  name: throw-dice-cron-job
spec:
  schedule: "30 21 * * *"
  jobTemplate:
    spec:
      completions: 3
      parallelism: 3
      backoffLimit: 25 # This is so the job does not quit before it succeeds.
      template:
        spec:
          containers:
            - name: throw-dice
              image: kodekloud/throw-dice
      restartPolicy: Never
```

44. Check the TargetPort configured in specific service svc

```
kubectl describe svc kubernetes | grep -i targetPort
```

```
controlplane ~ ➔ kubectl describe svc kubernetes | grep -i targetPort
TargetPort: 6443/TCP
```

45. Create a NodePort svc to access the web-app from outside of cluster

```
---
apiVersion: v1
kind: Service
metadata:
  name: webapp-service
spec:
  type: NodePort
  ports:
    - targetPort: 8080
      port: 8080
```

```
nodePort: 30080
selector:
  name: simple-webapp
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
service/kubernetes	ClusterIP	10.43.0.1	<none>	443/TCP	13m
service/webapp-service	NodePort	10.43.31.89	<none>	8080:30080/TCP	60s

46. Network Policy — *networkpolicy* — *netpol*

From internal to pod payroll traffic is allowed and any Internal pod access payroll pod using port 8080

```
kubectl get networkpolicy
kubectl describe networkpolicy payroll-policy
```

```
root@controlplane ~ ➔ kubectl get networkpolicy
NAME          POD-SELECTOR   AGE
payroll-policy  name=payroll  2m22s

root@controlplane ~ ➔ kubectl describe networkpolicy payroll-policy
Name:          payroll-policy
Namespace:    default
Created on:   2022-12-28 10:08:42 +0000 UTC
Labels:        <none>
Annotations:  <none>
Spec:
  PodSelector:  name=payroll
  Allowing ingress traffic:
    To Port: 8080/TCP
    From:
      PodSelector: name=internal
  Not affecting egress traffic
  Policy Types: Ingress
```

From internal to pod payroll traffic is allowed

47. Create Policy to allow pod access from outside

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: internal-policy
  namespace: default
spec:
  podSelector:
    matchLabels:
      name: internal
  policyTypes:
    - Egress
    - Ingress
  ingress:
    - {}
  egress:
    - to:
        - podSelector:
            matchLabels:
              name: mysql
  ports:
    - protocol: TCP
      port: 3306

    - to:
        - podSelector:
            matchLabels:
              name: payroll
  ports:
    - protocol: TCP
      port: 8080

  - ports:
      - port: 53
        protocol: UDP
      - port: 53
        protocol: TCP
kubectl describe netpol internal-policy
```

```
root@controlplane ~ ➔ kubectl describe netpol internal-policy
Name:           internal-policy
Namespace:      default
Created on:    2022-12-28 10:33:18 +0000 UTC
Labels:         <none>
Annotations:   <none>
Spec:
  PodSelector: name=internal
  Allowing ingress traffic:
    To Port: <any> (traffic allowed to all ports)
    From: <any> (traffic not restricted by source)
  Allowing egress traffic:
    To Port: 3306/TCP
    To:
      PodSelector: name=mysql
    -----
    To Port: 8080/TCP
    To:
      PodSelector: name=payroll
    -----
    To Port: 53/UDP
    To Port: 53/TCP
    To: <any> (traffic not restricted by destination)
Policy Types: Egress, Ingress
```

48. Check Ingress Resource in all namespace

```
kubectl get ingress --all-namespaces
kubectl describe ingress ingress-wear-watch -n app-space
```

```
controlplane ~ ➔ kubectl get ingress --all-namespaces
NAMESPACE  NAME          CLASS  HOSTS  ADDRESS        PORTS  AGE
app-space  ingress-wear-watch  <none> *  10.105.88.61  80     12m
```

```
controlplane ~ ➔ kubectl describe ingress ingress-wear-watch -n app-space
Name:          ingress-wear-watch
Labels:        <none>
Namespace:    app-space
Address:      10.105.88.61
Ingress Class: <none>
Default backend: <default>
Rules:
Host          Path  Backends
----          ----  -----
*
  /wear        wear-service:8080 (10.244.0.5:8080)
  /watch       video-service:8080 (10.244.0.4:8080)
Annotations:  nginx.ingress.kubernetes.io/rewrite-target: /
               nginx.ingress.kubernetes.io/ssl-redirect: false
Events:
Type  Reason  Age           From           Message
----  -----  --           --            --
Normal  Sync   14m (x2 over 14m)  nginx-ingress-controller  Scheduled for sync
```

49. Different Content is shown on /wear and /watch as per Backends configured in Ingress



50. Edit the key values in Ingress using edit

```
kubectl edit ingress ingress-wear-watch -n app-space
```

```
controlplane ~ ➔ kubectl edit ingress ingress-wear-watch -n app-space
ingress.networking.k8s.io/ingress-wear-watch edited
```

51. Create a new Ingress for path /pay and backend service name : pay-service

```
---
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: test-ingress
  namespace: critical-space
  annotations:
    nginx.ingress.kubernetes.io/rewrite-target: /
    nginx.ingress.kubernetes.io/ssl-redirect: "false"
spec:
  rules:
  - http:
    paths:
    - path: /pay
      pathType: Prefix
      backend:
        service:
          name: pay-service
          port:
            number: 8282
```

52. Roles and RoleBindings in Service Account

```
kubectl get roles -n ingress-space
kubectl describe roles -n ingress-space

kubectl get rolebindings -n ingress-space
kubectl describe rolebindings -n ingress-space
```

```
root@controlplane ~ ➔ kubectl get roles -n ingress-space
NAME          CREATED AT
ingress-role  2022-12-28T12:41:52Z

root@controlplane ~ ➔ kubectl describe roles -n ingress-space
Name:      ingress-role
Labels:    app.kubernetes.io/name=ingress-nginx
           app.kubernetes.io/part-of=ingress-nginx
Annotations: <none>
PolicyRule:
  Resources  Non-Resource URLs  Resource Names          Verbs
  -----      -----          -----
  configmaps []              []
  configmaps []              [ingress-controller-leader-nginx] [get update]
  endpoints   []              []
  namespaces  []              []
  pods        []              []
  secrets     []              []

root@controlplane ~ ➔ kubectl get rolebindings -n ingress-space
NAME          ROLE          AGE
ingress-role-binding  Role/ingress-role  4m6s

root@controlplane ~ ➔ kubectl describe rolebindings -n ingress-space
Name:      ingress-role-binding
Labels:    app.kubernetes.io/name=ingress-nginx
           app.kubernetes.io/part-of=ingress-nginx
Annotations: <none>
Role:
  Kind:  Role
  Name:  ingress-role
Subjects:
  Kind        Name          Namespace
  ----        ---          -----
  ServiceAccount  ingress-serviceaccount
```

53. Check Logs of Container from outside of Pod (From Host)

```
kubectl exec webapp -- cat /log/app.log
```

```
controlplane ~ ✘ kubectl exec webapp -- cat /log/app.log
[2022-12-28 14:27:28,268] INFO in event-simulator: USER3 is viewing page3
[2022-12-28 14:27:29,269] INFO in event-simulator: USER4 logged in
[2022-12-28 14:27:30,270] INFO in event-simulator: USER3 is viewing page2
[2022-12-28 14:27:31,271] INFO in event-simulator: USER4 is viewing page3
[2022-12-28 14:27:32,271] INFO in event-simulator: USER1 is viewing page2
[2022-12-28 14:27:32,271] WARNING in event-simulator: USER5 failed to log in
```

But If we delete the Container we won't be able to see logs on Host , but with volume we can do that .

***Create a volume “log-volume” with hostpath
“/var/log/webapp” and map it with /log in container
volumeMounts:***

```
apiVersion: v1
kind: Pod
metadata:
  name: webapp
spec:
  containers:
    - name: event-simulator
      image: kodekloud/event-simulator
      env:
        - name: LOG_HANDLERS
          value: file
      volumeMounts:
        - mountPath: /log
          name: log-volume

  volumes:
    - name: log-volume
      hostPath:
        # directory location on host
        path: /var/log/webapp
        # this field is optional
        type: Directory
```

54. Create a Persistent Volume

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: pv-log
spec:
  persistentVolumeReclaimPolicy: Retain
  accessModes:
    - ReadWriteMany
  capacity:
    storage: 100Mi
  hostPath:
    path: /pv/log
kubectl apply -f pv.yaml
kubectl get pv
kubectl describe pv
```

```
controlplane /var/log/webapp ➔ kubectl apply -f pv.yaml
persistentvolume/pv-log created

controlplane /var/log/webapp ➔ kubectl get pv
NAME      CAPACITY   ACCESS MODES   RECLAIM POLICY   STATUS
AGE
pv-log    100Mi      RWX           Retain          Available
8s

controlplane /var/log/webapp ➔ kubectl describe pv
Name:            pv-log
Labels:          <none>
Annotations:    <none>
Finalizers:     [kubernetes.io/pv-protection]
StorageClass:
Status:         Available
Claim:
Reclaim Policy: Retain
Access Modes:   RWX
VolumeMode:    Filesystem
Capacity:      100Mi
Node Affinity: <none>
Message:
Source:
  Type:        HostPath (bare host directory volume)
  Path:        /pv/log
  HostPathType:
Events:         <none>
```

55. Create a pvc Persistent Volume Claim

```
kind: PersistentVolumeClaim
apiVersion: v1
metadata:
  name: claim-log-1
spec:
  accessModes:
    - ReadWriteOnce
  resources:
```

```
requests:  
  storage: 50Mi  
kubectl apply -f pvc.yaml  
kubectl get pvc  
kubectl describe pvc claim-log-1
```

```
controlplane /var/log/webapp ➔ kubectl apply -f pvc.yaml  
persistentvolumeclaim/claim-log-1 created  
  
controlplane /var/log/webapp ➔ kubectl get pvc  
NAME      STATUS   VOLUME   CAPACITY   ACCESS MODES   STORAGECLASS   AGE  
claim-log-1  Pending  
  
controlplane /var/log/webapp ➔ kubectl describe pvc claim-log-1  
Name:           claim-log-1  
Namespace:      default  
StorageClass:  
Status:         Pending  
Volume:  
Labels:          <none>  
Annotations:     <none>  
Finalizers:     [kubernetes.io/pvc-protection]  
Capacity:  
Access Modes:  
VolumeMode:     Filesystem  
Used By:        <none>  
Events:  
  Type    Reason            Age            From                Message  
  ----  -----            ----          ----  
  Normal  FailedBinding  2s (x4 over 40s)  persistentvolume-controller  no persistent  
          available for this claim and no storage class is set
```

due to access mode mismatch pvc is in pending state

```
controlplane /var/log/webapp ➔ kubectl get pv,pvc  
NAME          CAPACITY   ACCESS MODES   RECLAIM POLICY   STATUS   CLAIM   STORAGE  
CLASS  REASON   AGE          RWX           Retain       Available  
persistentvolume/pv-log  100Mi       RWX           Retain       Available  
                         7m33s  
  
NAME          STATUS   VOLUME   CAPACITY   ACCESS MODES   STORAGECLASS   AGE  
AGE  
persistentvolumeclaim/claim-log-1  Pending  
3m35s
```

After changing access mode to ReadWriteMany pvc is in bound state

```
controlplane /var/log/webapp ➔ kubectl get pvc  
NAME      STATUS   VOLUME   CAPACITY   ACCESS MODES   STORAGECLASS   AGE  
claim-log-1  Bound    pv-log    100Mi     RWM          25s
```

Change the pvc from pv in pod yaml and delete/create pod

```
volumes:
- name: log-volume
  persistentVolumeClaim:
    claimName: claim-log-1
```

56. STORAGE CLASS

```
---
apiVersion: storage.k8s.io/v1
kind: StorageClass
metadata:
  name: delayed-volume-sc
provisioner: kubernetes.io/no-provisioner
volumeBindingMode: WaitForFirstConsumer
kubectl get sc
kubectl describe sc local-path
```

```
controlplane ~ ➔ kubectl get sc
NAME          PROVISIONER           RECLAIMPOLICY   VOLUMEBINDINGMODE   ALLOWVOLUM
EEXPANSION   AGE
local-path (default)  rancher.io/local-path  Delete        WaitForFirstConsumer  false
            3m28s

controlplane ~ ➔ kubectl describe sc local-path
Name:           local-path
IsDefaultClass: Yes
Annotations:   objectset.rio.cattle.io/applied=H4sIAAAAAAAA/4yRT+vUMBCGv4rMuab1bu1KwIOu7
EUEQdDzNJlux6aZkkwry7LfXbIqrIffn2PyZN7hfXIFXPg7xcQSwEBSiXimaupSxfJ2q6GAiYMDA9/+oKPH1KCAmRQdKoK5A
oYgisoSUj5K/50sJtIqs1QWVT3lNM4xUDzJ5VegWJ63CQxMTXogW128+czBvf/gnIQXIwLOBAa8WPTl30qvGkoL2jw5rT2V6
ZKUZj+SbG5eZVRDKR0F8SpdDTg6rW8YzCgcSW4FeCxJ/+sjxHTCAbqrhmag20Pw9DbZtfu210z7JuhpNQ719m2w3c0e7fPc
f81W1DHfL1E2Th/IEUwEDHYkWJe8PCsgJgl8PxVPNsLGPhEnjRr2cSvM33k4Dicv4jLC34g60niiWPSo4S0zhTh9jsAAP//y
tgh5S0CAA,objectset.rio.cattle.io/id=,objectset.rio.cattle.io/owner-gvk=k3s.cattle.io/v1, Kind=
Addon,objectset.rio.cattle.io/owner-name=local-storage,objectset.rio.cattle.io/owner-namespace=k
ube-system,storageclass.kubernetes.io/is-default-class=true
Provisioner:   rancher.io/local-path
Parameters:    <none>
AllowVolumeExpansion: <unset>
MountOptions:  <none>
ReclaimPolicy: Delete
VolumeBindingMode: WaitForFirstConsumer
Events:       <none>
```

57. create a Pod using persistentVolumeClaim

```
---  
apiVersion: v1  
kind: Pod  
metadata:  
  name: nginx  
  labels:  
    name: nginx  
spec:  
  containers:  
    - name: nginx  
      image: nginx:alpine  
      volumeMounts:  
        - name: local-persistent-storage  
          mountPath: /var/www/html  
  volumes:  
    - name: local-persistent-storage  
      persistentVolumeClaim:  
        claimName: local-pvc
```

58. switch to other context using kubectl

```
kubectl config --kubeconfig=/root/my-kube-config use-context research
```

```
controlplane ~ ✘ kubectl config --kubeconfig=/root/my-kube-config use-context research  
Switched to context "research".
```

59. make the new my-kube-config as default config file

```
cp my-kube-config .kube/config
```

```
=====
```

60. Identify the authorization modes configured on the cluster.

```
kubectl get pods -n kube-system
kubectl describe pod kube-apiserver-controlplane -n kube-system | grep -i
authorization
```

```
controlplane ~ ➔ kubectl get pods -n kube-system
NAME                               READY   STATUS    RESTARTS   AGE
coredns-6d4b75cb6d-nnnls          1/1     Running   0          9m47s
coredns-6d4b75cb6d-v27mc          1/1     Running   0          9m47s
etcd-controlplane                 1/1     Running   0          9m58s
kube-apiserver-controlplane       1/1     Running   0          10m
kube-controller-manager-controlplane 1/1     Running   0          9m58s
kube-flannel-ds-2mmqz            1/1     Running   0          9m47s
kube-proxy-dnwzk                 1/1     Running   0          9m47s
kube-scheduler-controlplane       1/1     Running   0          10m
controlplane ~ ➔ kubectl describe pod kube-apiserver-controlplane -n kube-system | grep -i authorization
--authorization-mode=Node,RBAC
```

61. Roles

Check the access to resources under a specific Role

```
controlplane ~ ➔ kubectl get roles --all-namespaces
NAMESPACE      NAME                           CREATED AT
blue           developer                      2022-12-29T06:57:44Z
kube-public    kubeadm:bootstrap-signer-clusterinfo 2022-12-29T06:49:19Z
kube-public    system:controller:bootstrap-signer   2022-12-29T06:49:17Z
kube-system    extension-apiserver-authentication-reader 2022-12-29T06:49:17Z
kube-system    kube-proxy                      2022-12-29T06:49:19Z
kube-system    kubeadm:kubelet-config             2022-12-29T06:49:17Z
kube-system    kubeadm:nodes-kubeadm-config        2022-12-29T06:49:17Z
kube-system    system::leader-locking-kube-controller-manager 2022-12-29T06:49:17Z
kube-system    system::leader-locking-kube-scheduler 2022-12-29T06:49:17Z
kube-system    system:controller:bootstrap-signer   2022-12-29T06:49:17Z
kube-system    system:controller:cloud-provider       2022-12-29T06:49:17Z
kube-system    system:controller:token-cleaner       2022-12-29T06:49:17Z
```

```
controlplane ~ ➔ kubectl describe role kube-proxy -n kube-system
Name:          kube-proxy
Labels:        <none>
Annotations:   <none>
PolicyRule:
  Resources      Non-Resource URLs  Resource Names  Verbs
  -----          -----           -----           -----
  configmaps     []               [kube-proxy]    [get]
```

kube-proxy Role can only get/view details of configmaps

62. Which account is mapped to the Role kube-proxy

```
kubectl get rolebindings -n kube-system
kubectl describe rolebinding kube-proxy -n kube-system
```

```
controlplane ~ ➔ kubectl get rolebindings -n kube-system
NAME                           ROLE                                         AGE
kube-proxy                      Role/kube-proxy                         27m
kubeadm:kubelet-config          Role/kubeadm:kubelet-config          27m
kubeadm:nodes-kubeadm-config   Role/kubeadm:nodes-kubeadm-config   27m
system::extension-apiserver-authentication-reader
system::leader-locking-kube-controller-manager
system::leader-locking-kube-scheduler
system:controller:bootstrap-signer
system:controller:cloud-provider
system:controller:token-cleaner

controlplane ~ ➔ kubectl describe rolebinding kube-proxy -n kube-system
Name:          kube-proxy
Labels:        <none>
Annotations:   <none>
Role:
  Kind:  Role
  Name:  kube-proxy
Subjects:
  Kind  Name
  ----  ---
  Group system:bootstrappers:kubeadm:default-node-token
```

63. execute command using other user - --as

```
kubectl get pods --as dev-user
```

```
controlplane ~ ➔ kubectl get pods --as dev-user
Error from server (Forbidden): pods is forbidden: User "dev-user" cannot list resource "pods" in API group "" in
the namespace "default"
```

64. Create a Role and roleBinding to access Pods from dev-user

```
kind: Role
apiVersion: rbac.authorization.k8s.io/v1
metadata:
  namespace: default
  name: developer
rules:
- apiGroups: [""]
  resources: ["pods"]
  verbs: ["list", "create", "delete"]

---
kind: RoleBinding
apiVersion: rbac.authorization.k8s.io/v1
metadata:
  name: dev-user-binding
subjects:
- kind: User
  name: dev-user
  apiGroup: rbac.authorization.k8s.io
roleRef:
  kind: Role
  name: developer
  apiGroup: rbac.authorization.k8s.io
```

```
controlplane ~ ➔ kubectl apply -f role_roleBinding.yaml
role.rbac.authorization.k8s.io/developer created
rolebinding.rbac.authorization.k8s.io/dev-user-binding created
```

65. Update the Role to create Pod using deployment .

```
kind: Role
apiVersion: rbac.authorization.k8s.io/v1
metadata:
  namespace: blue
  name: developer
rules:
- apiGroups: ["apps"]
  resources: ["pods"]
  resourceNames: ["darf-blue-app"]
  verbs: ["list", "create", "delete"]
- apiGroups: ["apps"]
  resources: ["deployments"]
  verbs: ["list", "create", "delete", "watch"]
```

66. Count the number of clusterroles in cluster using wc

```
kubectl get clusterroles --all-namespaces | wc -l
```

```
controlplane ~ ➔ kubectl get clusterroles --all-namespaces | wc -l
70
```

ClusterRoles are cluster wide and not part of any namespace

```
=====
=====
```

67. Add storage access to existing role michelle

```
kubectl api-resources | grep -i storageclass
kubectl api-resources | grep -i persistentvol
```

```
controlplane ~ ➔ kubectl api-resources | grep -i storageclass
storageclasses           sc          storage.k8s.io/v1                   false      StorageClass
controlplane ~ ➔ kubectl api-resources | grep -i persistentvol
persistentvolumeclaims   pvc         v1                           true       PersistentVolumeClaim
persistentvolumes        pv          v1                           false      PersistentVolume
```

```
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
  name: michelle-role
rules:
- apiGroups: ["v1"]
  resources: ["nodes"]
  verbs: ["get", "list", "watch", "delete", "create"]
- apiGroups: ["storage.k8s.io/v1"]
  resources: ["storageclasses"]
  verbs: ["get", "list", "watch", "delete", "create"]
- apiGroups: ["v1"]
  resources: ["persistentvolumes"]
  verbs: ["get", "list", "watch", "delete", "create"]
---
kind: ClusterRoleBinding
apiVersion: rbac.authorization.k8s.io/v1
metadata:
  name: michelle-binding
subjects:
- kind: User
  name: michelle
  apiGroup: rbac.authorization.k8s.io
roleRef:
  kind: ClusterRole
  name: michelle-role
  apiGroup: rbac.authorization.k8s.io
```

68. Add *NamespaceAutoProvision* in — *enable-admission-plugins* to allow creation of namespace while creating Pod in new namespace

vi /etc/kubernetes/manifests/kube-apiserver.yaml

```
apiVersion: v1
kind: Pod
metadata:
  annotations:
    kubeadm.kubernetes.io/kube-apiserver.advertise-add
  creationTimestamp: null
  labels:
    component: kube-apiserver
    tier: control-plane
  name: kube-apiserver
  namespace: kube-system
spec:
  containers:
    - command:
      - kube-apiserver
      - --advertise-address=10.32.39.9
      - --allow-privileged=true
      - --authorization-mode=Node,RBAC
      - --client-ca-file=/etc/kubernetes/pki/ca.crt
      - --enable-admission-plugins=NodeRestriction
      - --enable-bootstrap-token-auth=true
```

After changes

```
containers:
  - command:
    - kube-apiserver
    - --advertise-address=10.32.39.9
    - --allow-privileged=true
    - --authorization-mode=Node,RBAC
    - --client-ca-file=/etc/kubernetes/pki/ca.crt
    - --enable-admission-plugins=NodeRestriction,NamespaceAutoProvision
    - --enable-bootstrap-token-auth=true
    - --etcd-cafile=/etc/kubernetes/pki/etcd/ca.crt
    - --etcd-certfile=/etc/kubernetes/pki/apiserver-etcd-client.crt
```



`NamespaceExists` and `NamespaceAutoProvision` admission controllers are deprecated and now replaced by `NamespaceLifecycle` admission controller.

The `NamespaceLifecycle` admission controller will make sure that requests to a non-existent namespace is rejected and that the default namespaces such as `default`, `kube-system` and `kube-public` cannot be deleted.

=====

69. Create a TLS secret for webhook

```
kubectl -n webhook-demo create secret tls webhook-server-tls \
--cert "/root/keys/webhook-server-tls.crt" \
--key "/root/keys/webhook-server-tls.key"
```

70. Find short names for multiple resources in api-resources

```
kubectl api-resources | grep -E -i
'deployments|replicasets|cronjobs|customresourcedefinitions'
```

```
controlplane ~ ➔ kubectl api-resources | grep -E -i 'deployments|replicasets|cronjobs|customresourcedefinitions'
customresourcedefinitions      crd,crds   apiextensions.k8s.io/v1           false      CustomResourceDefinition
deployments                   deploy     apps/v1                         true       Deployment
replicasets                  rs        apps/v1                         true       ReplicaSet
cronjobs                      cj        batch/v1                        true       CronJob
```

71. Find major minor in kubernetes version

```
kubectl version -o yaml | grep -i -E 'major|minor'
```

```
controlplane ~ → kubectl version -o yaml | grep -i -E 'major|minor'  
  major: "1"  
  minor: "24"  
  major: "1"  
  minor: "24"
```

72. Enable the `v1alpha1` version for `rbac.authorization.k8s.io` API group

Take a backup of that `apiserver` manifest file

```
cp -v /etc/kubernetes/manifests/kube-apiserver.yaml /root/kube-  
apiserver.yaml.backup
```

Edit the `apiserver` manifest file :

```
vi /etc/kubernetes/manifests/kube-apiserver.yaml
```

Add below line :

```
- --runtime-config=rbac.authorization.k8s.io/v1alpha1
```

kubelet will recreate `apiserver` pod in sometime

```
controlplane ~ ✘ kubectl get po -n kube-system
The connection to the server controlplane:6443 was refused - did you specify the
right host or port?

controlplane ~ ✘ kubectl get po -n kube-system
NAME                      READY   STATUS    RESTARTS   AGE
coredns-6d4b75cb6d-kpb7z   1/1     Running   0          39m
coredns-6d4b75cb6d-qqfv8   1/1     Running   0          39m
etcd-controlplane          1/1     Running   0          39m
kube-apiserver-controlplane 0/1     Pending   0          1s
kube-controller-manager-controlplane 1/1     Running   1 (28s ago) 39m
kube-flannel-ds-cmv9s      1/1     Running   0          39m
kube-proxy-9bwv7            1/1     Running   0          39m
kube-scheduler-controlplane 1/1     Running   1 (28s ago) 39m
```

=====

73. change api version using kubectl-convert plugin

```
# Install plugin
curl -LO https://dl.k8s.io/release/$(curl -L -s
https://dl.k8s.io/release/stable.txt)/bin/linux/amd64/kubectl-convert
chmod +x kubectl-convert
mv kubectl-convert /usr/local/bin/kubectl-convert
kubectl-convert -h

# convert api version
kubectl-convert -f ingress-old.yaml --output-version networking.k8s.io/v1 >
ingress-new.yaml
kubectl create -f ingress-new.yaml
kubectl get ing ingress-space -oyaml | grep apiVersion
```

74. custom Resource (*customresourcedefinitions*)

It is an extension of the Kubernetes API that is not necessarily available in a default Kubernetes installation , It is either namespace or cluster scope based .

```
kind: CustomResourceDefinition
metadata:
  name: internals.datasets.kodecloud.com
spec:
  group: datasets.kodecloud.com
  versions:
    - name: v1
      served: true
      storage: true
      schema:
        openAPIV3Schema:
          type: object
          properties:
            spec:
              type: object
              properties:
                internalLoad:
                  type: string
                range:
                  type: integer
                percentage:
                  type: string
  scope: namespaced
  names:
    plural: internal
    singular: internal
    kind: Internal
    shortNames:
      - int
vi crd.yaml
kubectl apply -f crd.yaml
kubectl get customresourcedefinitions
kubectl describe customresourcedefinitions internals.datasets.kodekloud.com
```

```
controlplane ~ ➔ kubectl apply -f crd.yaml
customresourcedefinition.apirextensions.k8s.io/internals.datasets.kodekloud.com created
```

```
controlplane ~ ➔ kubectl get customresourcedefinitions
NAME                                CREATED AT
collectors.monitoring.controller    2022-12-30T04:57:23Z
globals.traffic.controller          2022-12-30T04:57:24Z
internals.datasets.kodekloud.com    2022-12-30T05:15:47Z

controlplane ~ ➔ kubectl describe customresourcedefinitions internals.datasets.kodekloud.com
Name:           internals.datasets.kodekloud.com
Namespace:
Labels:         <none>
Annotations:   <none>
API Version:  apiextensions.k8s.io/v1
Kind:          CustomResourceDefinition
Metadata:
  Creation Timestamp: 2022-12-30T05:15:47Z
  Generation:        1
  Managed Fields:
```

75. Create a new custom resource using customresourcedefinitions

```
---
kind: Internal
apiVersion: datasets.kodekloud.com/v1
metadata:
  name: internal-space
  namespace: default
spec:
  internalLoad: "high"
  range: 80
  percentage: "50"
~
vi custom.yaml
kubectl apply -f custom.yaml
kubectl api-resources | grep -i Kodekloud
kubectl get internals
kubectl describe internals internal-space
```

```
controlplane ~ ➔ kubectl apply -f custom.yaml
internal.datasets.kodekloud.com/internal-space created

controlplane ~ ➔ kubectl api-resources | grep -i Kodekloud
internals          int      datasets.kodekloud.com/v1          true
  Internal
```

```
controlplane ~ ➔ kubectl get internals
NAME          AGE
internal-space 3m17s

controlplane ~ ➔ kubectl describe internals internal-space
Name:           internal-space
Namespace:      default
Labels:         <none>
Annotations:   <none>
API Version:  datasets.kodekloud.com/v1
Kind:          Internal
Metadata:
  Creation Timestamp: 2022-12-30T05:21:23Z
  Generation:        1
```

=====

====

76. Deployment Strategy

Canary

Recreate

Blue-Green

Rolling Update

```
kubectl describe deployments.apps frontend | grep -i StrategyType:
```

```
controlplane ~ ➔ kubectl describe deployments.apps frontend | grep -i StrategyType:
StrategyType:          RollingUpdate
```

```
=====
```

====

77. decode password from secret using base64 -d

```
echo Password: $(kubectl get secret --namespace default bravo-drupal -o jsonpath=".data.drupal-password" | base64 -d)  
kubectl describe secret bravo-drupal
```

```
controlplane ~ ➜ echo Password: $(kubectl get secret --namespace default bravo-drupal -o jsonpath=".data.drupal-password" | base64 -d)  
Password: UKL1Hz5PN  
  
controlplane ~ ➜ kubectl describe secret bravo-drupal  
Name:           bravo-drupal  
Namespace:      default  
Labels:         app.kubernetes.io/instance=bravo  
                app.kubernetes.io/managed-by=Helm  
                app.kubernetes.io/name=drupal  
                helm.sh/chart=drupal-12.5.12  
Annotations:    meta.helm.sh/release-name: bravo  
                meta.helm.sh/release-namespace: default  
  
Type:          Opaque  
  
Data  
=====  
drupal-password: 10 bytes
```

78. upgrading from Client side to server side and vice versa

```
kubectl apply --server-side [--dry-run=server]
```

By default, field management of the object transfers from client-side apply to kubectl server-side apply without encountering conflicts.

Downgrading

```
kubectl apply --server-side --field-manager=my-manager [--dry-run=server]
```

Thanks .



[Ranjeet Jangra](#)

Network Automation Expert with 13 years of experience in Development | Testing | Deployment | Support | Automation on various Technologies like IP-Routing, Cloud, Programming, Containers, Kubernetes, Telemetry, Orchestration, Network-Programmability, YANG, TextFSM, Jinja, RestAPI

<https://www.linkedin.com/in/ranjeetjangra/>